

Experiment 2: Loan Amount Prediction using Linear Regression

Shinigdapriya Sathish

August 16, 2025

Aim

To predict the loan amount sanctioned to users using Linear Regression on a real-world dataset and interpret the results through proper evaluation and visualization.

Libraries Used

- pandas
- numpy
- matplotlib
- seaborn
- scikit-learn

Objective

Apply Linear Regression to predict the sanctioned loan amount, evaluate the model using error metrics and cross-validation, and visualize the results to gain insights into model performance.

Mathematical Background

The hypothesis function for Linear Regression is:

$$\hat{y} = w_0 + w_1x_1 + w_2x_2 + \cdots + w_nx_n$$

The model minimizes the Mean Squared Error (MSE):

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

The coefficients are learned using the Normal Equation:

$$w = (X^T X)^{-1} X^T y$$

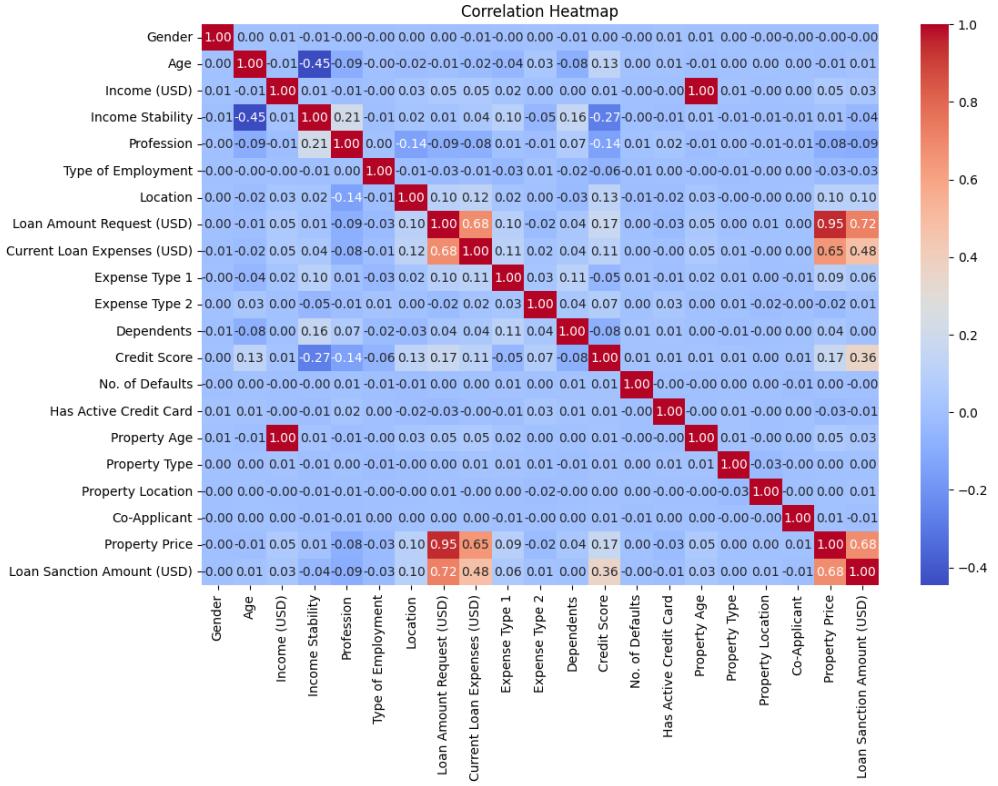
Preprocessing

- Dropped irrelevant columns (e.g., Customer ID, Name, Property ID)
- Filled missing values (mode for categorical, median for numeric)
- Encoded categorical variables using Label Encoding
- Standardized numerical features using StandardScaler

Exploratory Data Analysis

Visualizations include:

- Histogram and Boxplot of `Income (USD)` and `Loan Sanction Amount (USD)`
- Correlation heatmap of numerical features
- Scatter plots for feature-target relationships
- Residual plots to assess linearity assumptions



Model Training

Used LinearRegression from scikit-learn. The dataset was split into 80% training and 20% validation sets.

Evaluation Metrics

- Mean Absolute Error (MAE)
- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- R² Score
- Adjusted R² Score

```
[ ]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
```

```

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

```

[]: train_path = '/content/drive/MyDrive/Docs/archive/train.csv'
df = pd.read_csv(train_path)

[]: # Step 1: Drop irrelevant columns
df.drop(columns=['Customer ID', 'Name', 'Property ID'], inplace=True)

[]: # Step 2: Handle missing values (simple strategy: fill with mode for categorical, median for numeric)
for col in df.columns:
 if df[col].dtype == 'object':
 df[col] = df[col].fillna(df[col].mode()[0])
 else:
 df[col] = df[col].fillna(df[col].median())

Step 3: Encode categorical columns
label_encoders = {}
for col in df.select_dtypes(include='object').columns:
 le = LabelEncoder()
 df[col] = le.fit_transform(df[col])
 label_encoders[col] = le

Step 4: Feature / Target split
target = 'Loan Sanction Amount (USD)'
X = df.drop(columns=[target])
y = df[target]

Step 5: Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

Step 6: Train-validation split
X_train, X_val, y_train, y_val = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

Step 7: Train model
lr = LinearRegression()
lr.fit(X_train, y_train)

Step 8: Evaluate

```
y_pred = lr.predict(X_val)

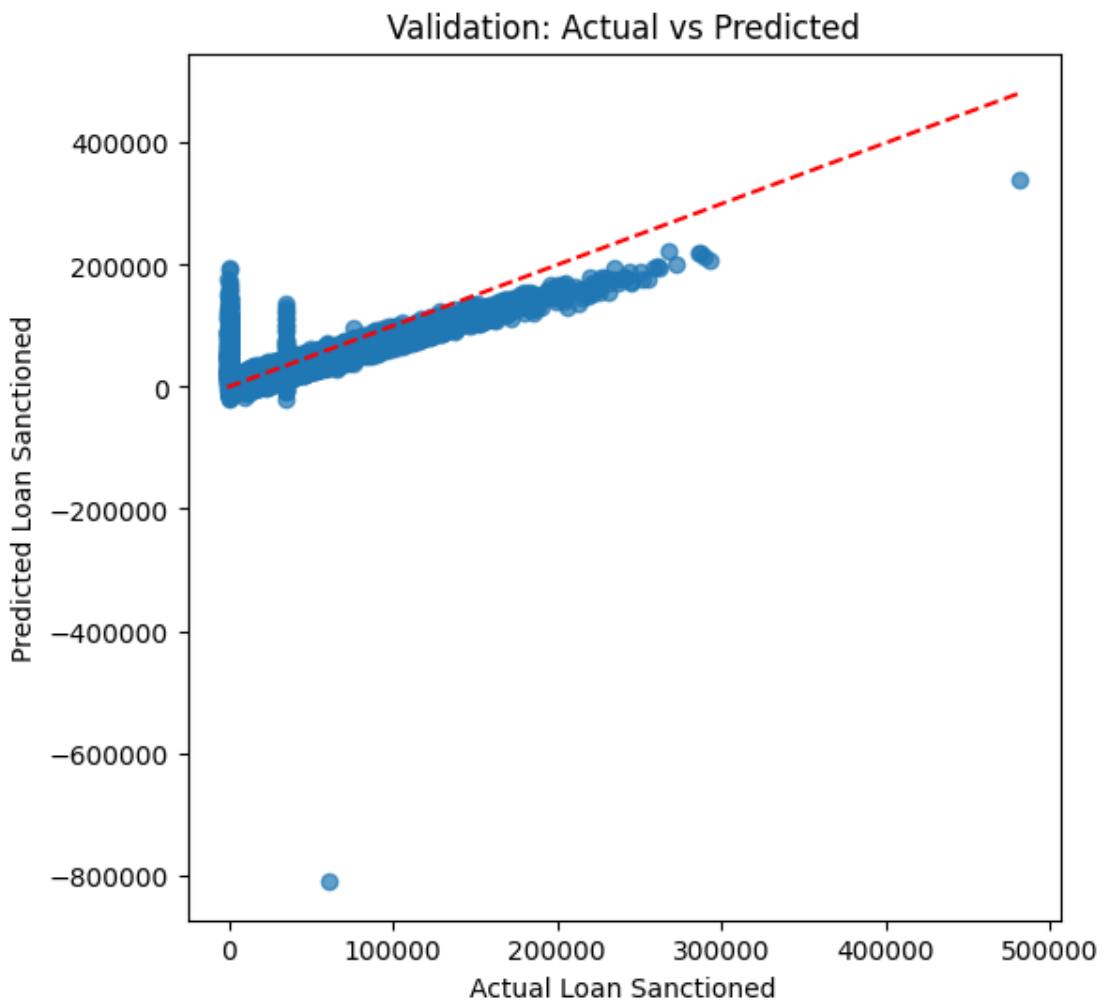
print("MAE:", mean_absolute_error(y_val, y_pred))
print("MSE:", mean_squared_error(y_val, y_pred))
print("R2 Score:", r2_score(y_val, y_pred))

# Step 9: Visualization
plt.figure(figsize=(6,6))
plt.scatter(y_val, y_pred, alpha=0.7)
plt.xlabel("Actual Loan Sanctioned")
plt.ylabel("Predicted Loan Sanctioned")
plt.title("Validation: Actual vs Predicted")
plt.plot([y_val.min(), y_val.max()], [y_val.min(), y_val.max()], 'r--')
plt.show()
```

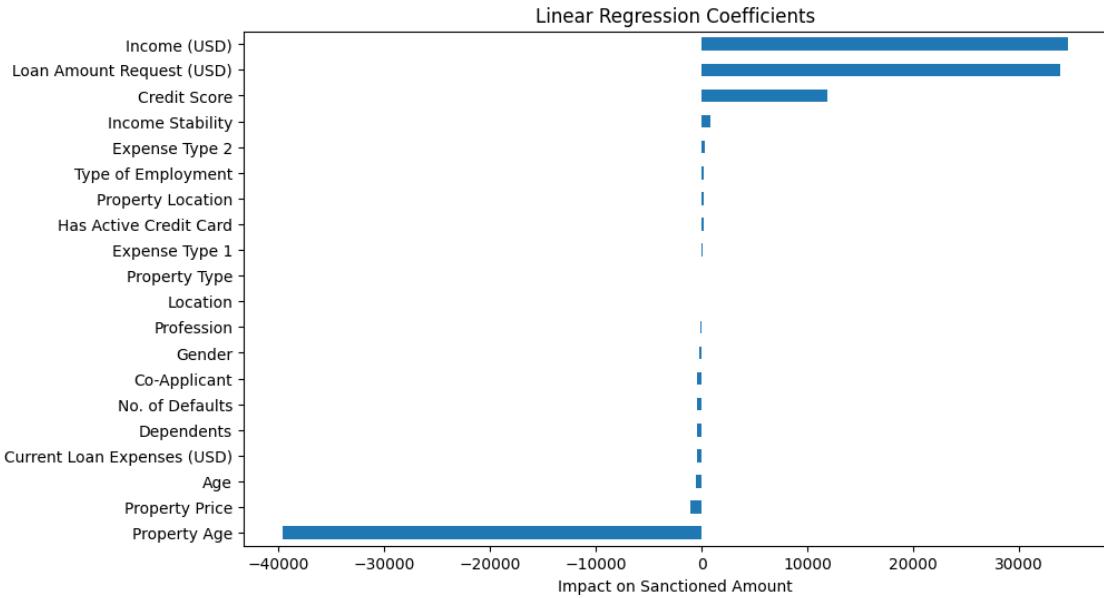
MAE: 21799.224445462914

MSE: 1076701697.7600782

R² Score: 0.531982511846756



```
[ ]: # Feature importance
coef = pd.Series(lr.coef_, index=X.columns)
coef.sort_values().plot(kind='barh', figsize=(10,6), title="Linear Regression Coefficients")
plt.xlabel("Impact on Sanctioned Amount")
plt.show()
```



```
[ ]: import seaborn as sns
import matplotlib.pyplot as plt

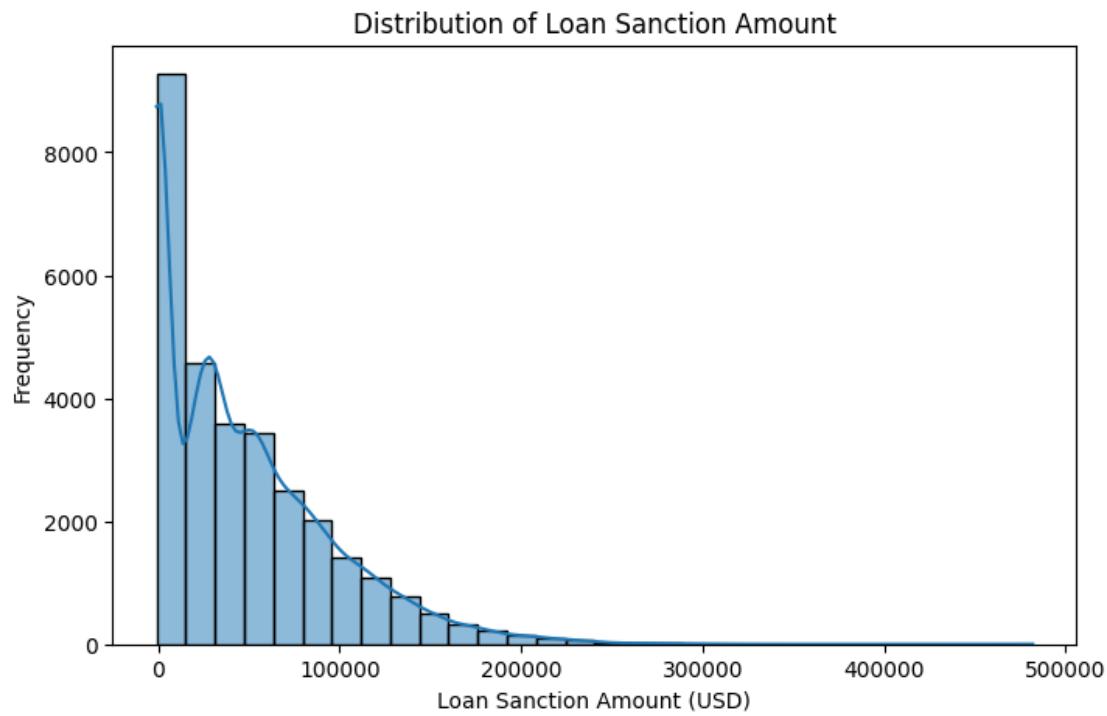
# Distribution of Loan Sanction Amount
plt.figure(figsize=(8, 5))
sns.histplot(df['Loan Sanction Amount (USD)'], kde=True, bins=30)
plt.title('Distribution of Loan Sanction Amount')
plt.xlabel('Loan Sanction Amount (USD)')
plt.ylabel('Frequency')
plt.show()

plt.figure(figsize=(8, 5))

numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns
print(numeric_cols)

for col in numeric_cols:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[col], kde=True, bins=30, color='teal')
    plt.title(f'Distribution of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')
    plt.grid(True)
    plt.tight_layout()
```

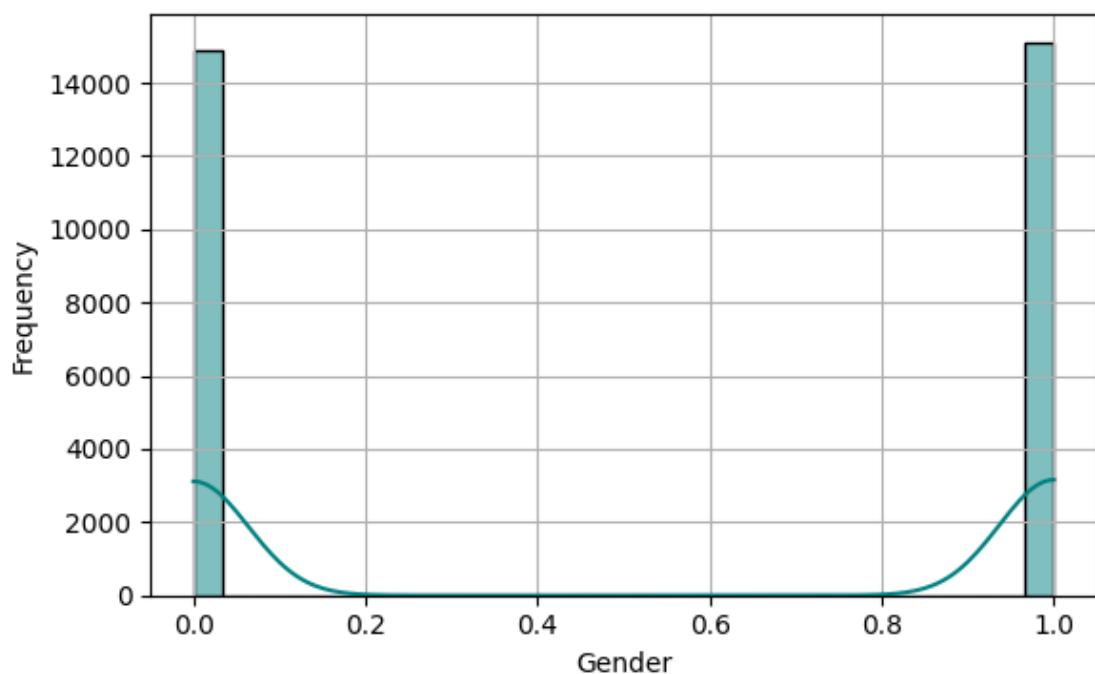
```
plt.show()
```



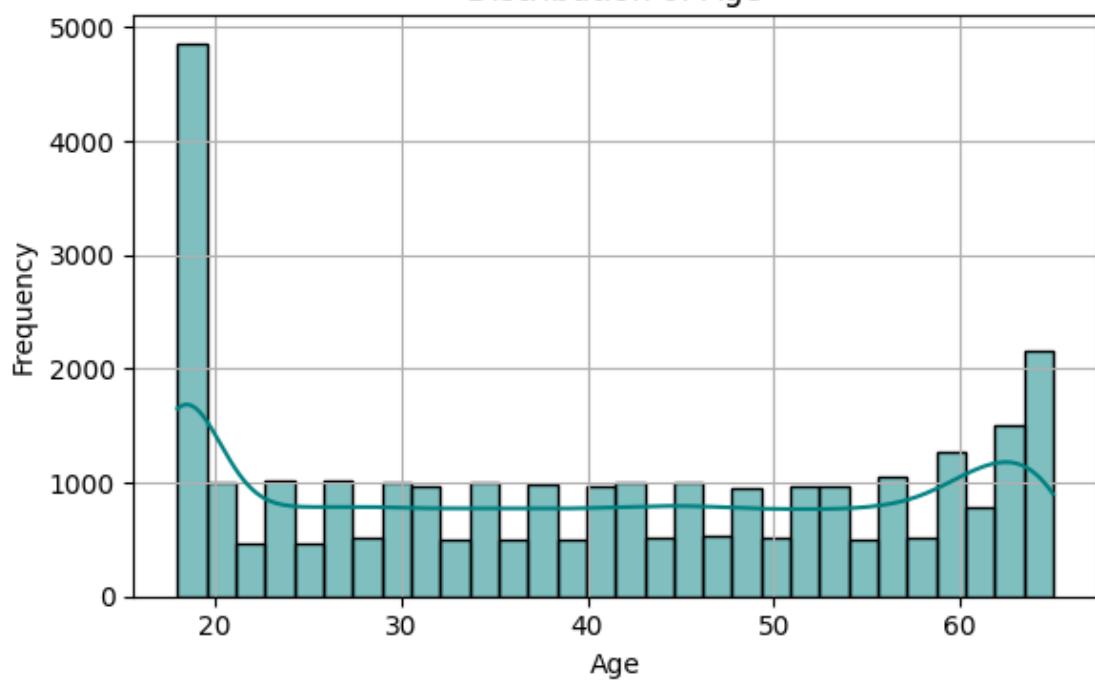
```
Index(['Gender', 'Age', 'Income (USD)', 'Income Stability', 'Profession',
       'Type of Employment', 'Location', 'Loan Amount Request (USD)',
       'Current Loan Expenses (USD)', 'Expense Type 1', 'Expense Type 2',
       'Dependents', 'Credit Score', 'No. of Defaults',
       'Has Active Credit Card', 'Property Age', 'Property Type',
       'Property Location', 'Co-Applicant', 'Property Price',
       'Loan Sanction Amount (USD)'],
      dtype='object')
```

<Figure size 800x500 with 0 Axes>

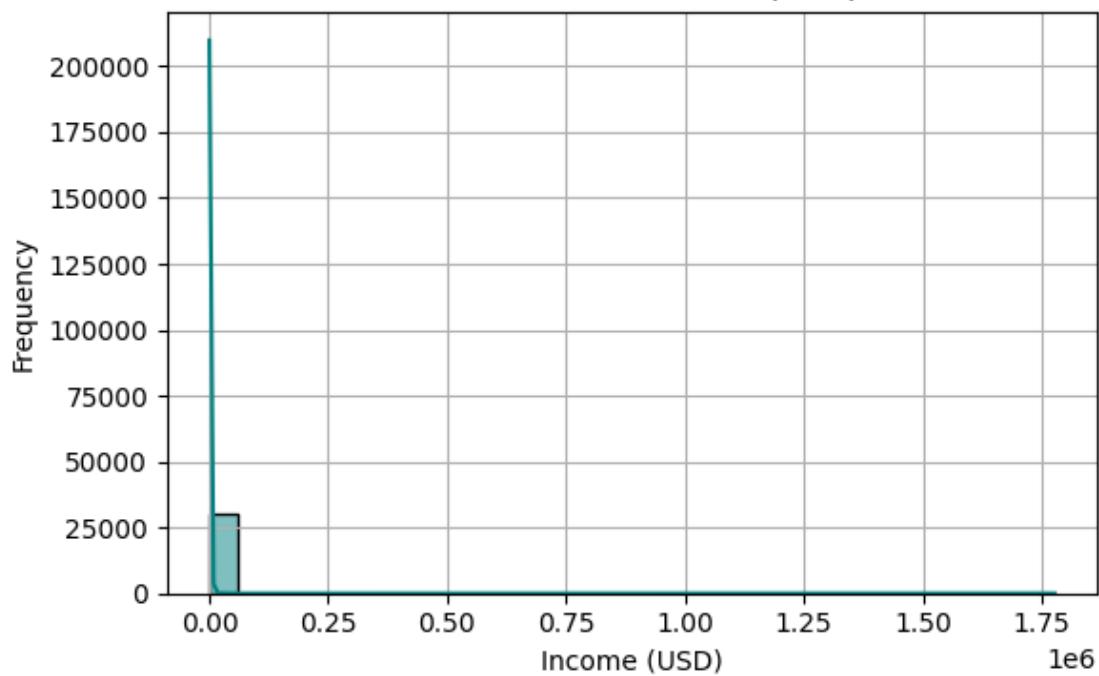
Distribution of Gender



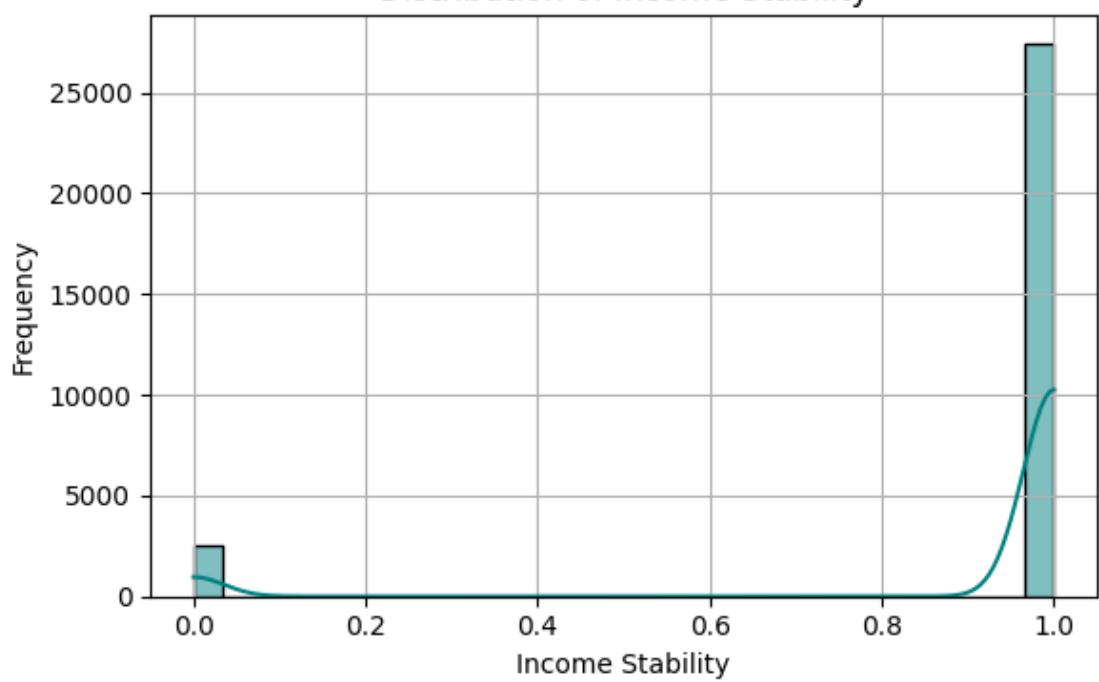
Distribution of Age



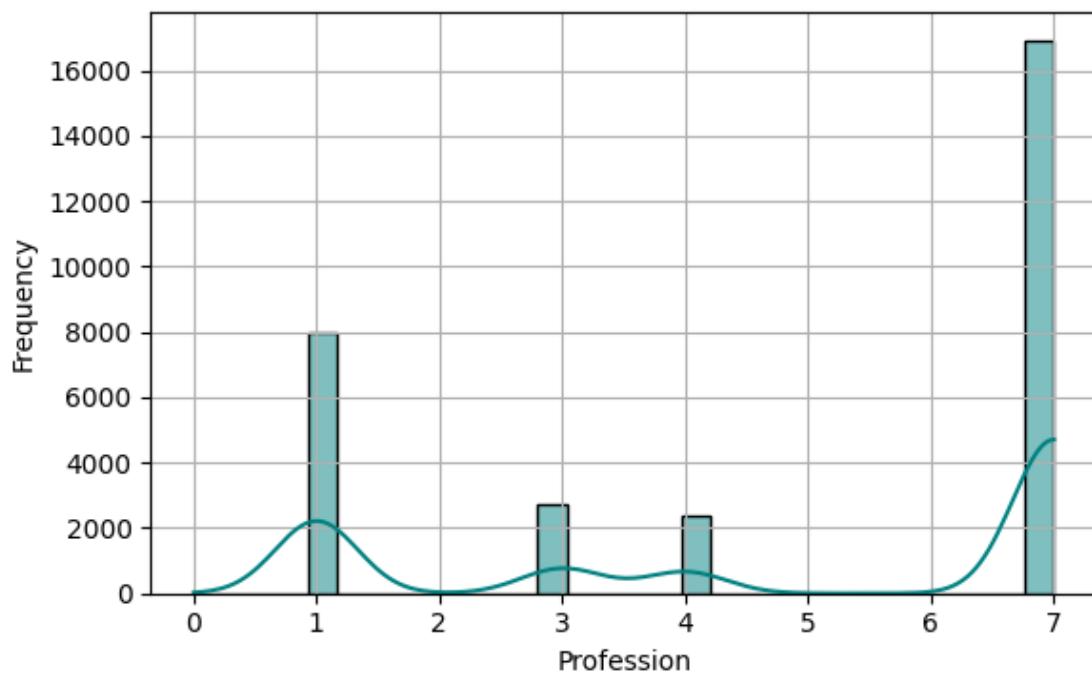
Distribution of Income (USD)



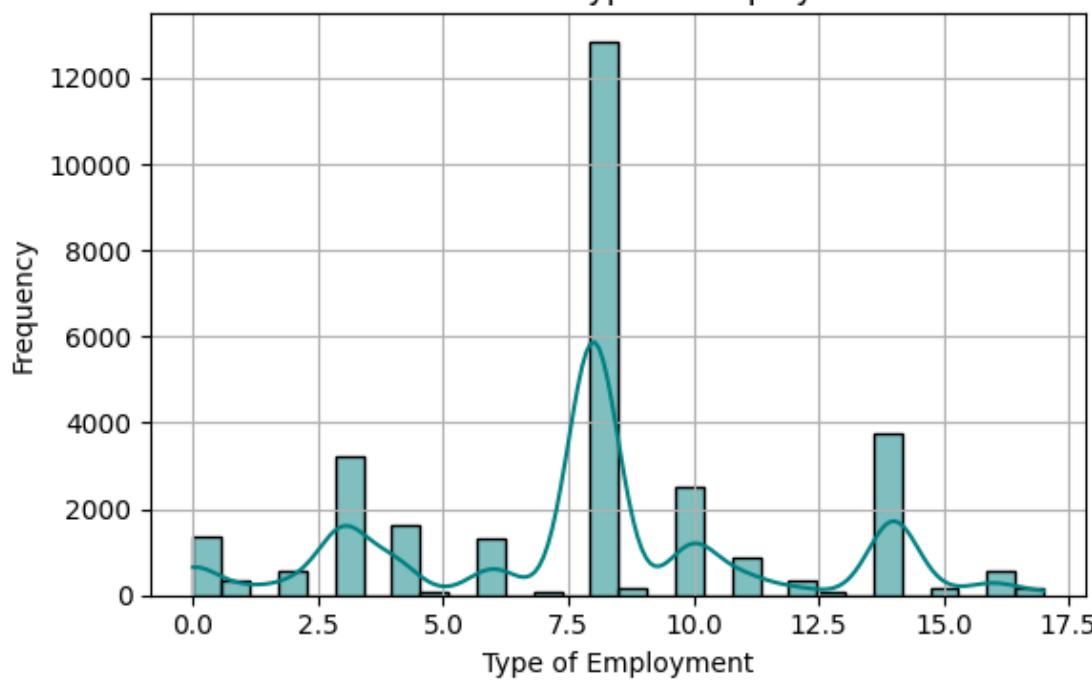
Distribution of Income Stability



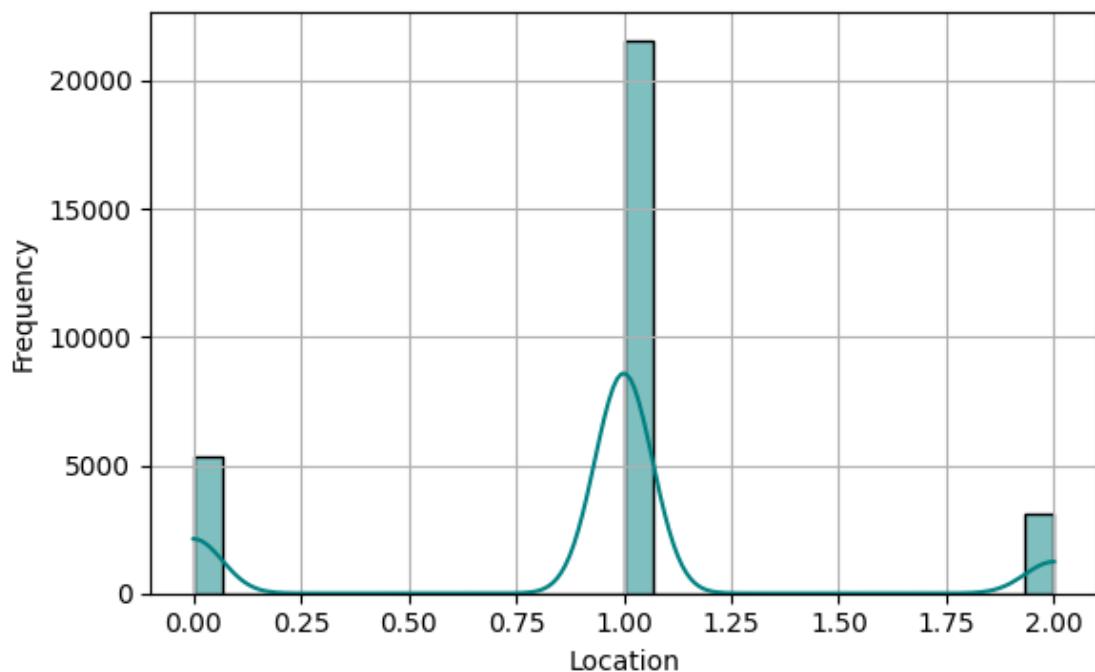
Distribution of Profession



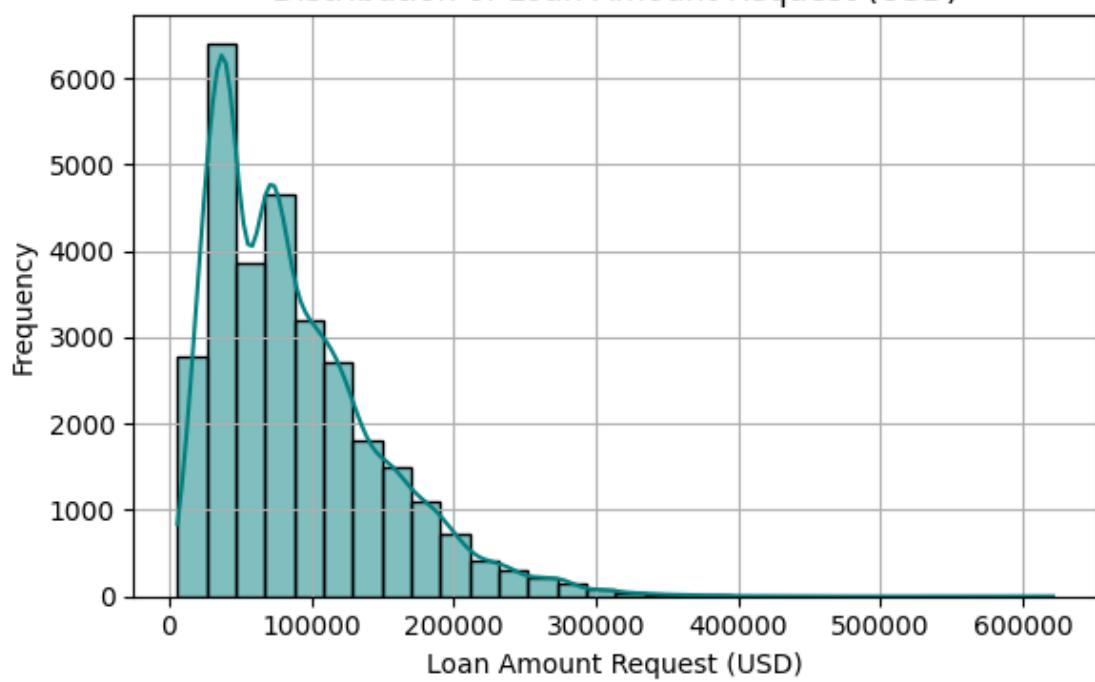
Distribution of Type of Employment



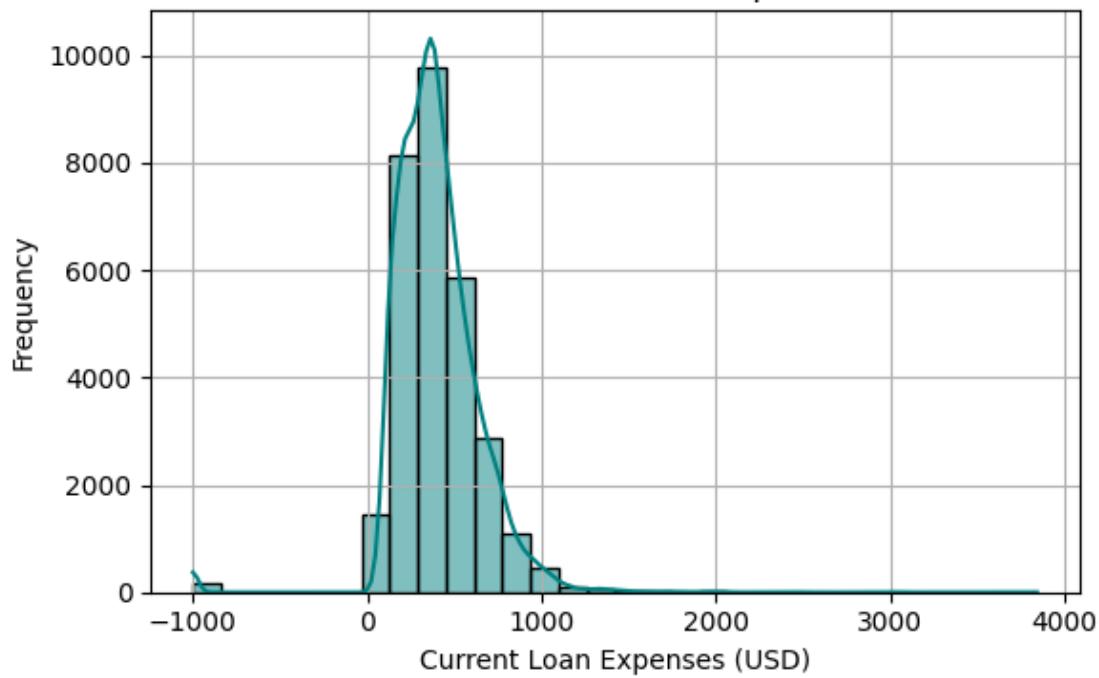
Distribution of Location



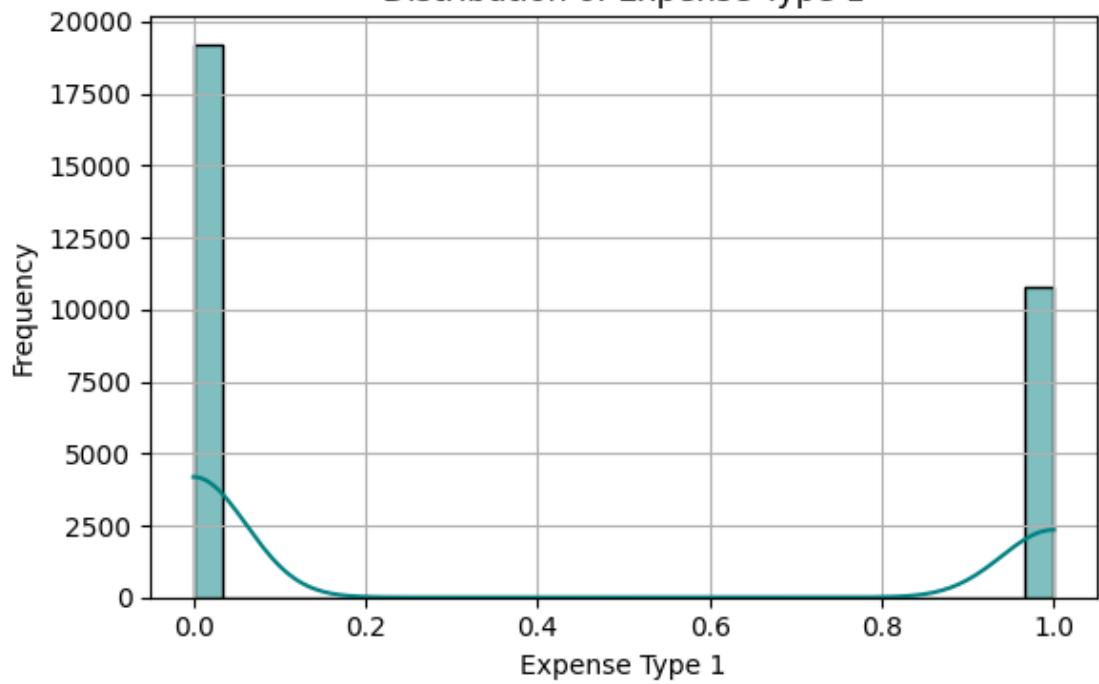
Distribution of Loan Amount Request (USD)



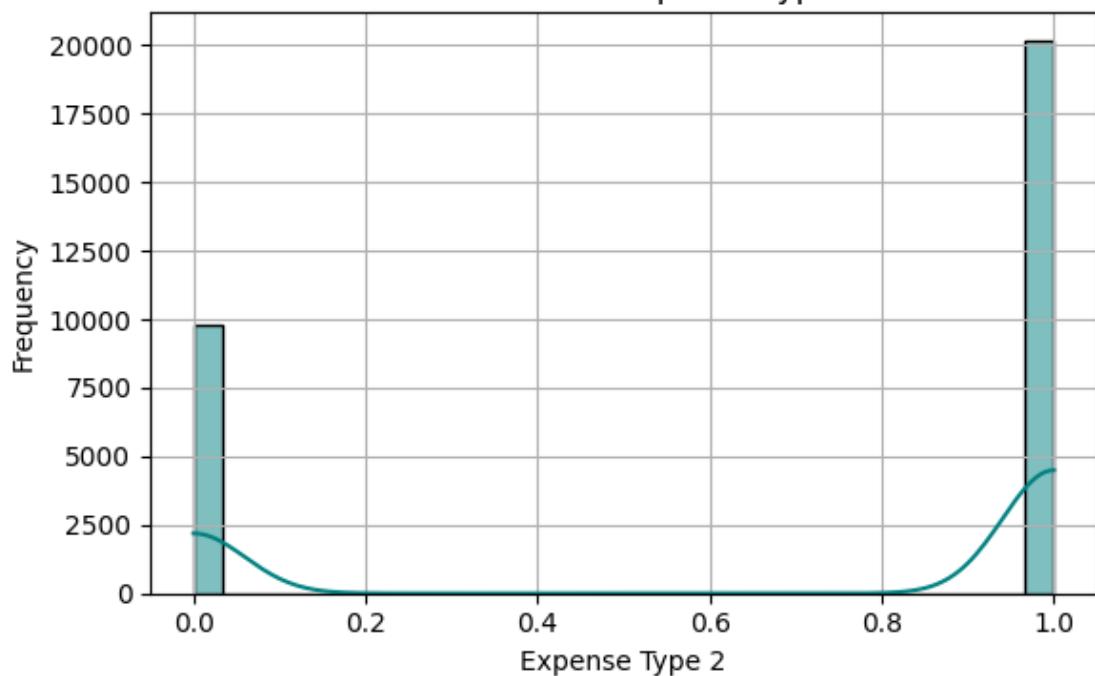
Distribution of Current Loan Expenses (USD)



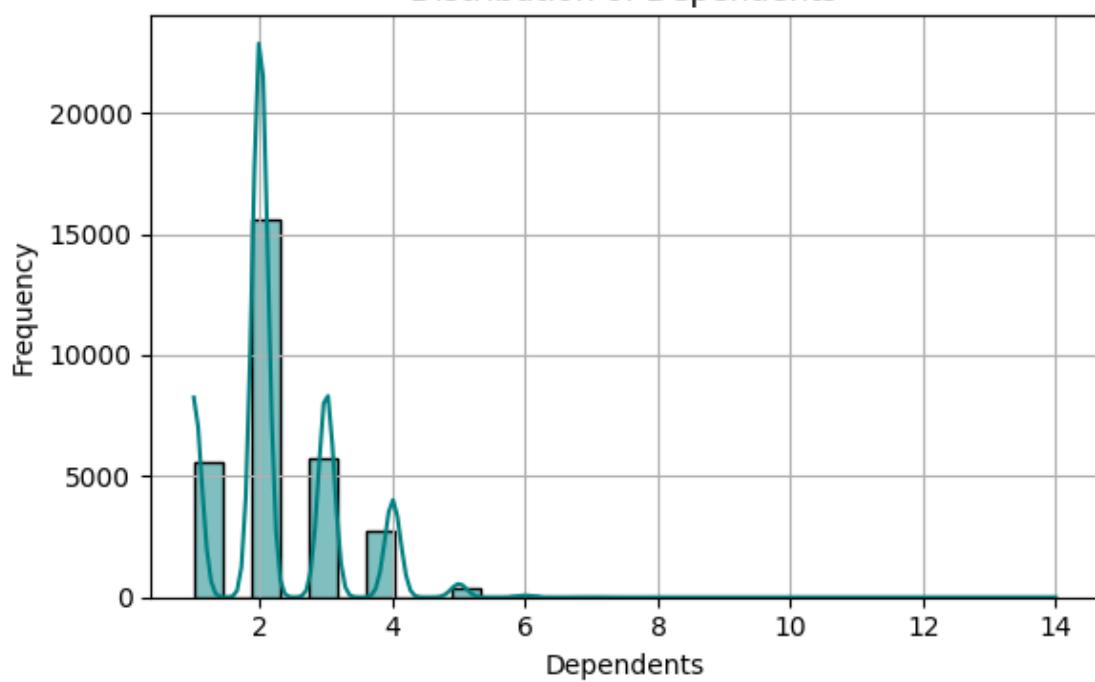
Distribution of Expense Type 1

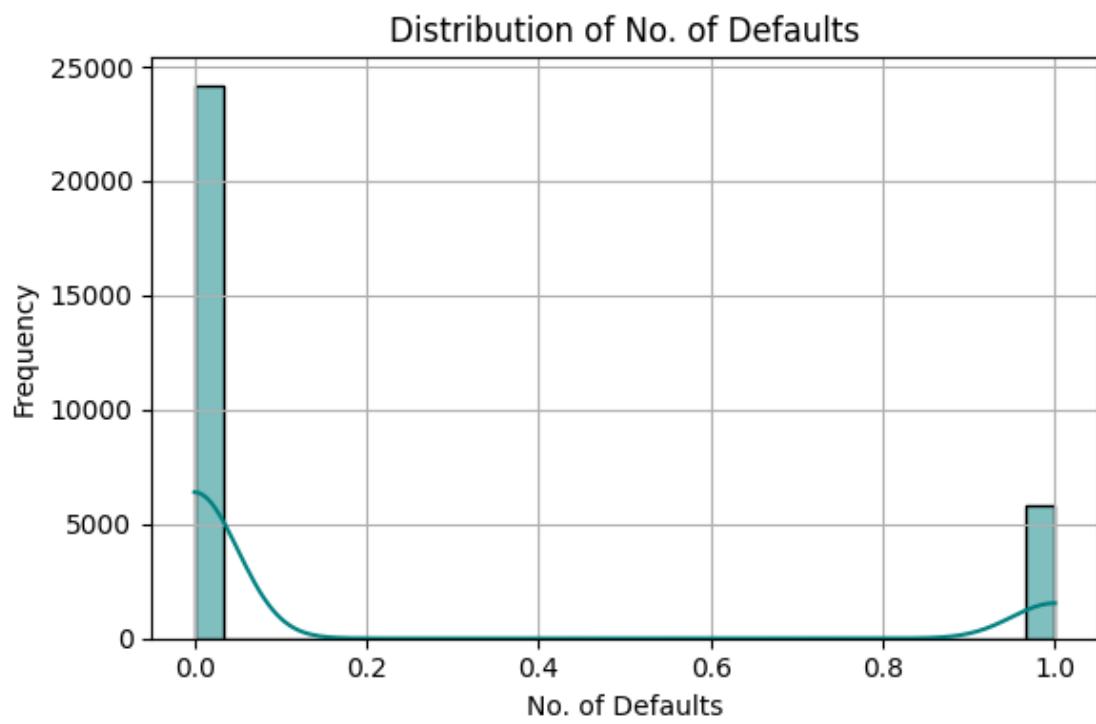
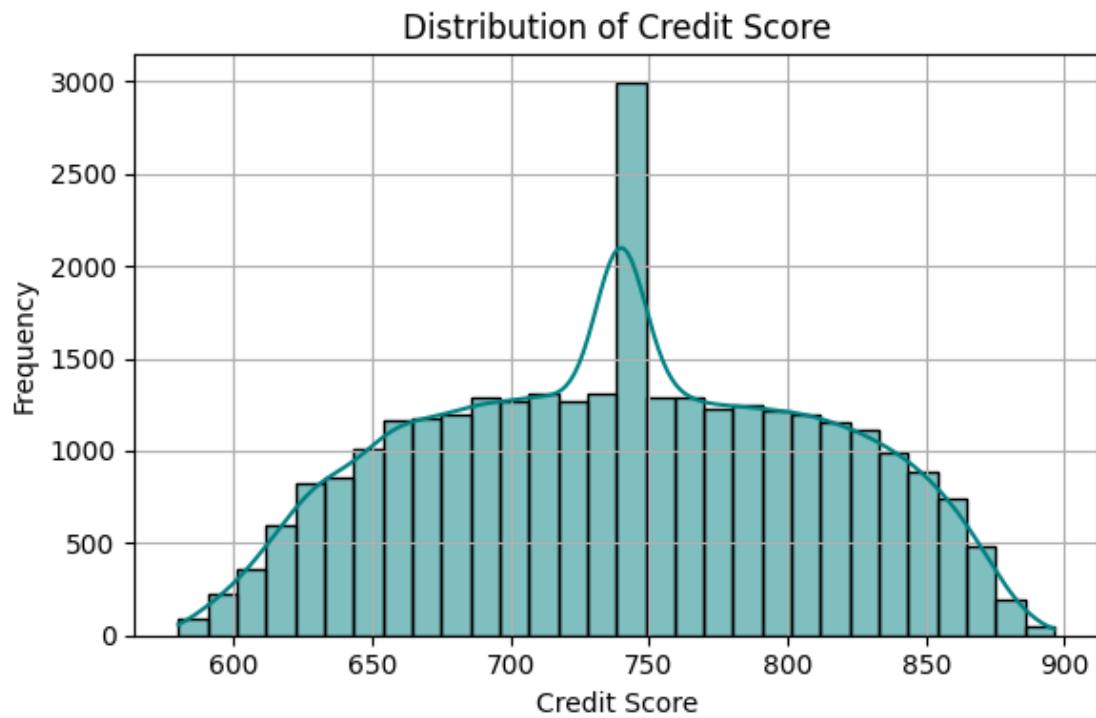


Distribution of Expense Type 2

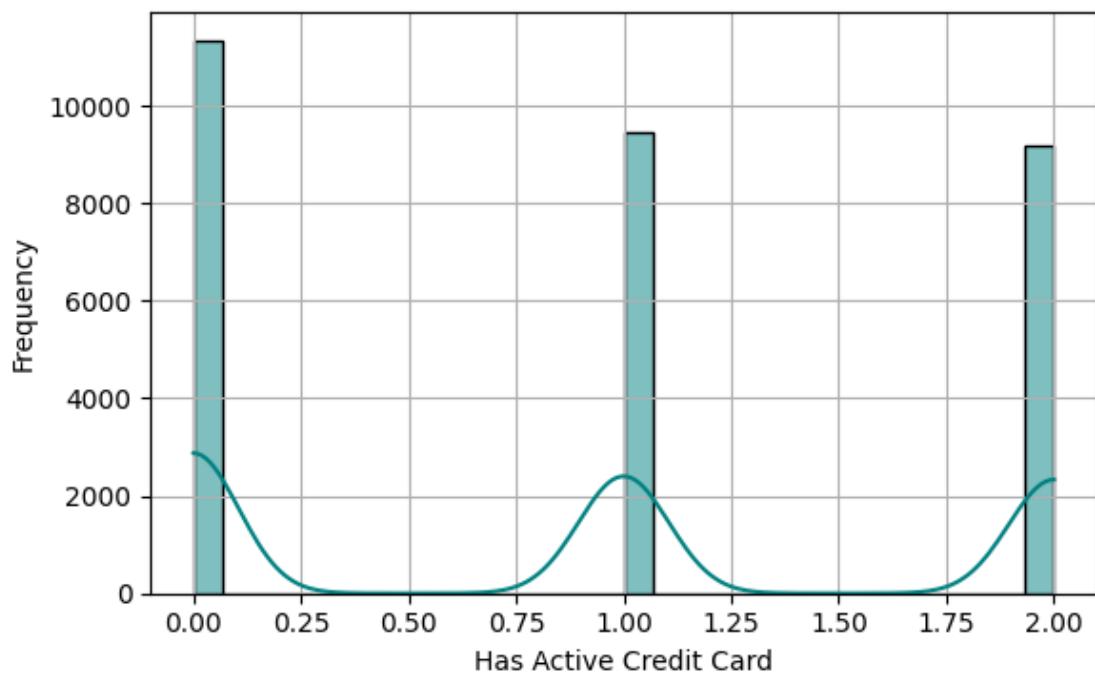


Distribution of Dependents

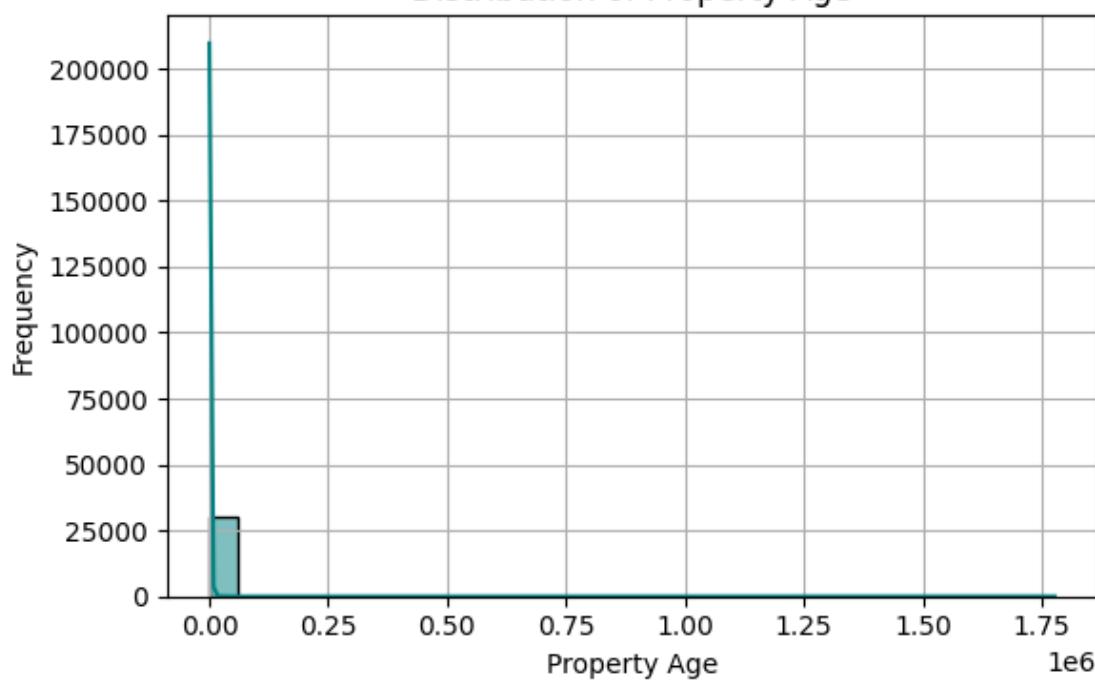




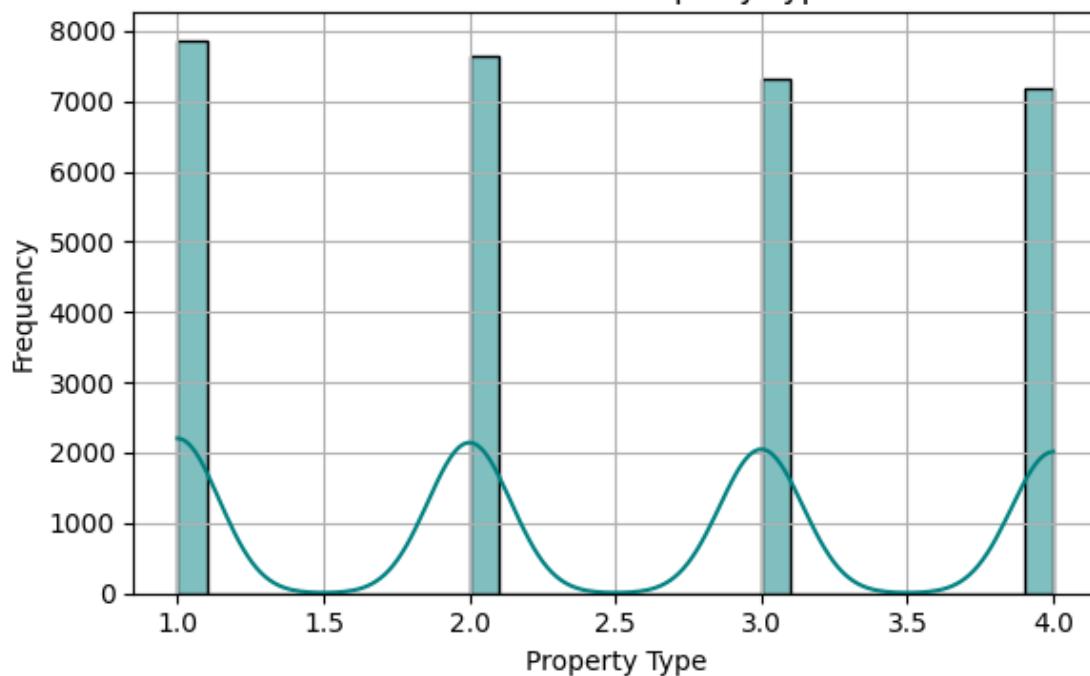
Distribution of Has Active Credit Card



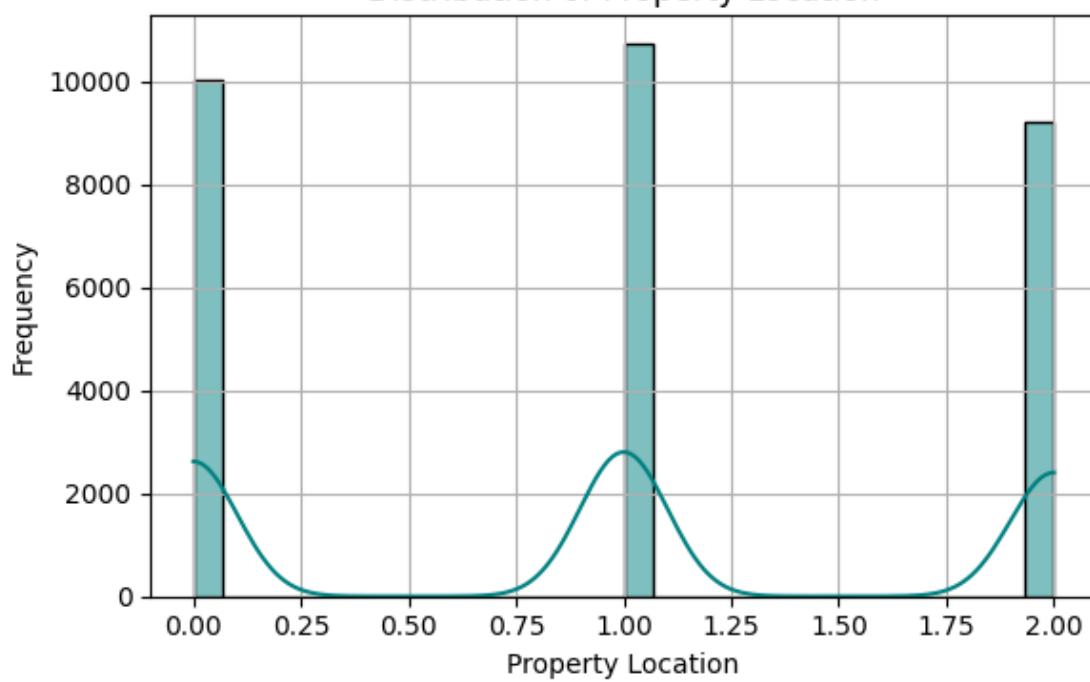
Distribution of Property Age

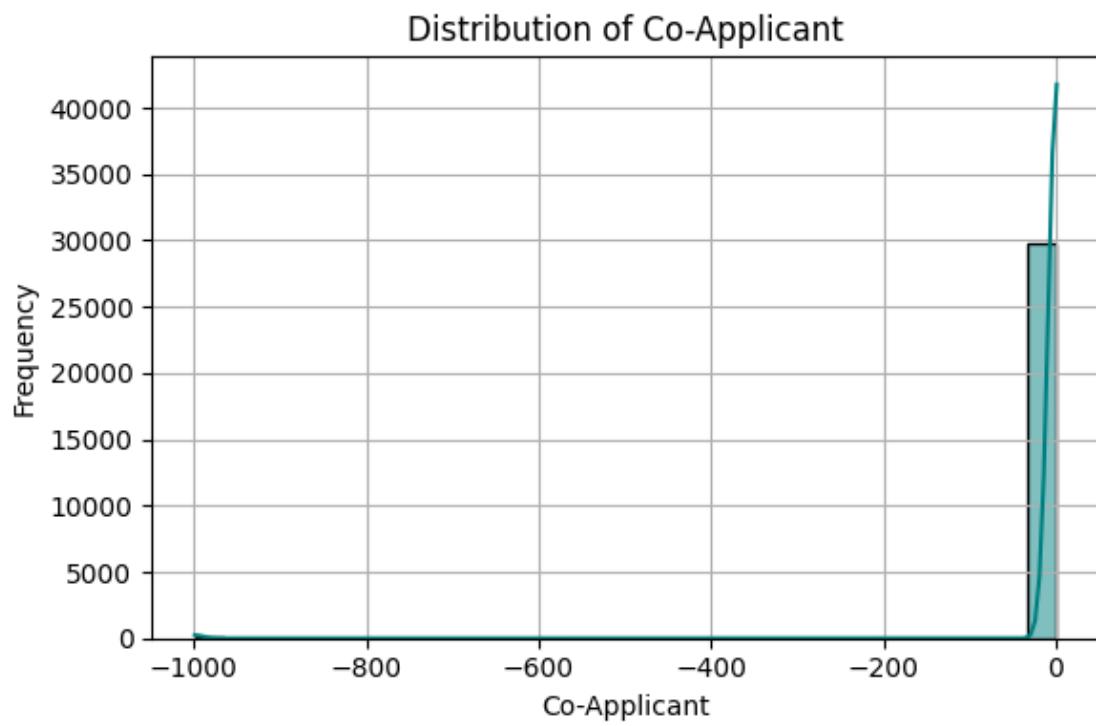


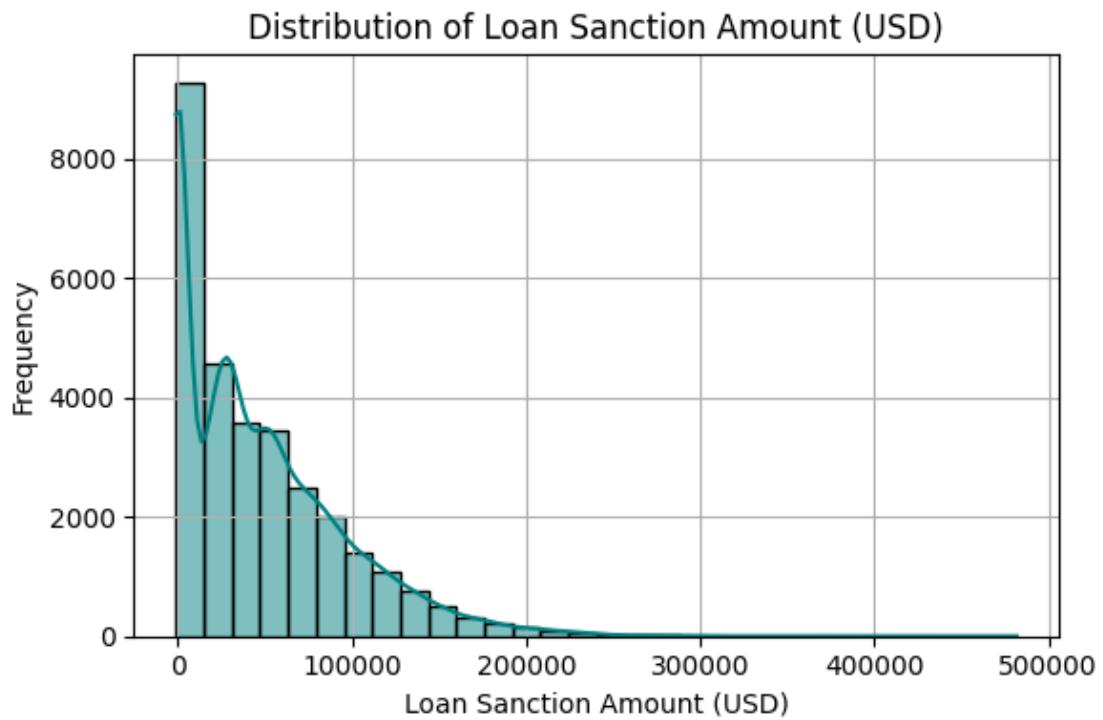
Distribution of Property Type



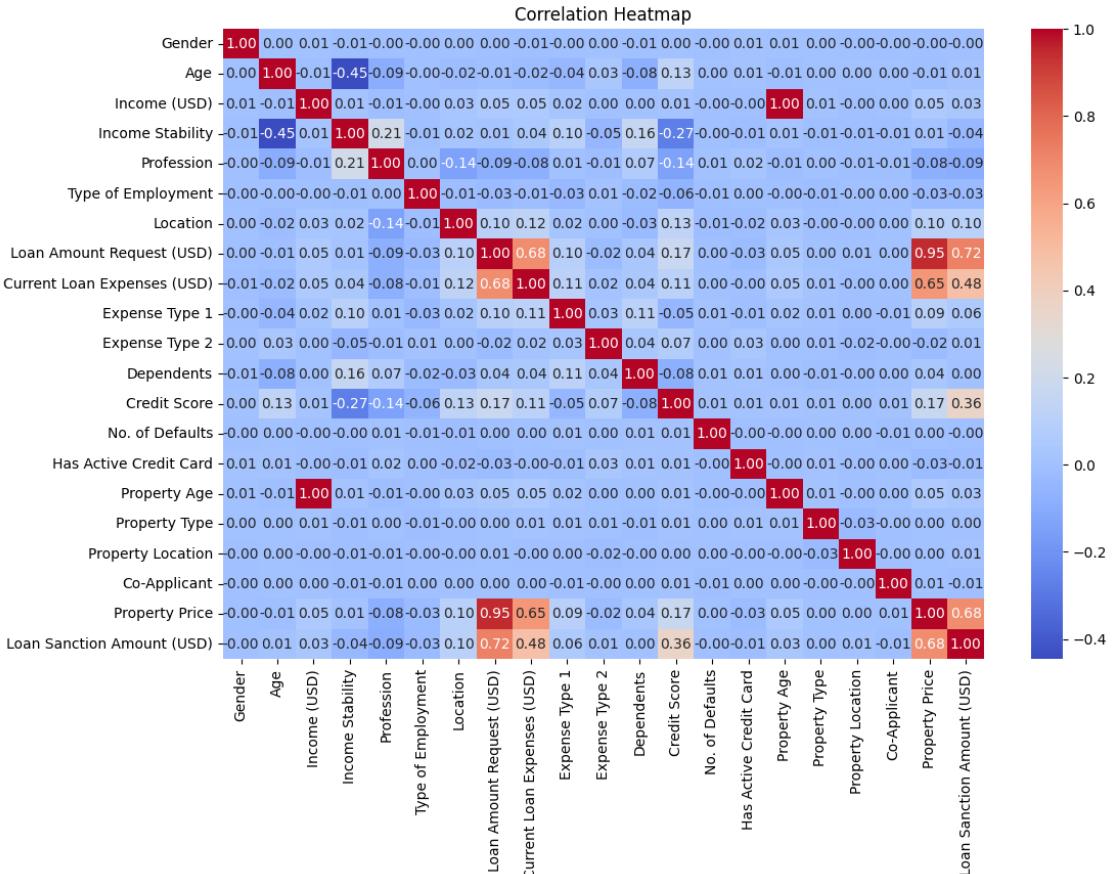
Distribution of Property Location



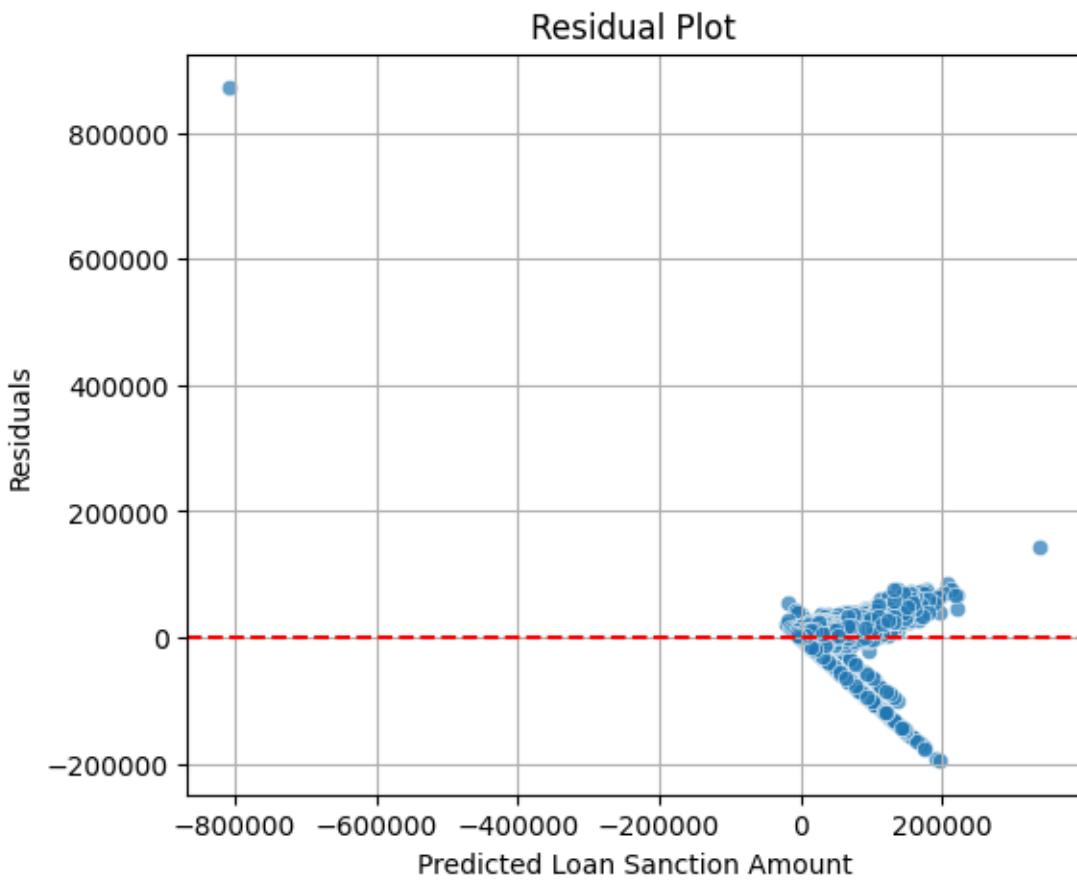




```
[ ]: plt.figure(figsize=(12, 8))
corr_matrix = df.corr()
sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```



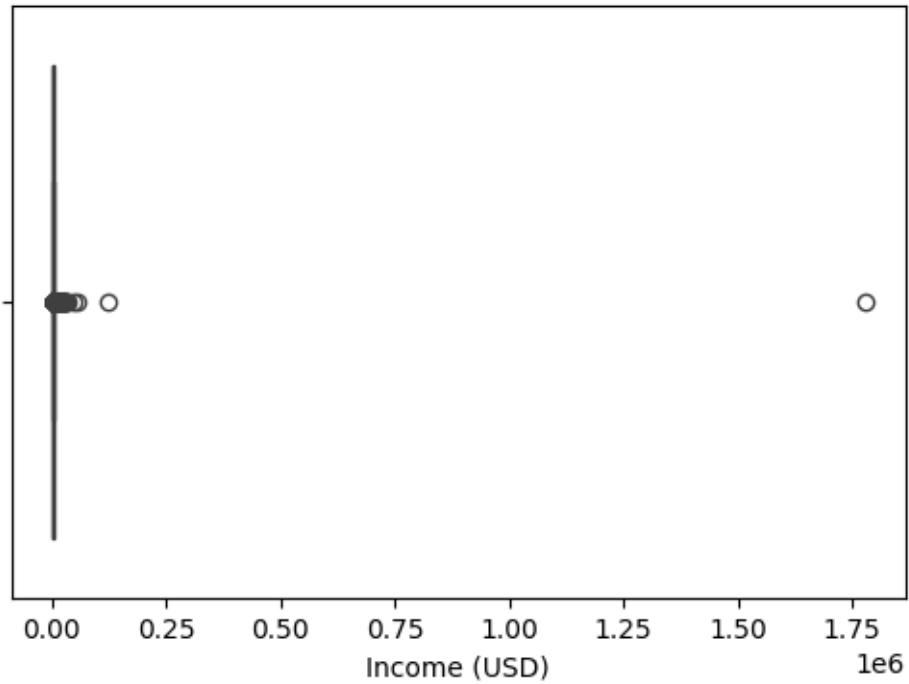
```
[ ]: residuals = y_val - y_pred
plt.figure(figsize=(6, 5))
sns.scatterplot(x=y_pred, y=residuals, alpha=0.7)
plt.axhline(0, color='red', linestyle='--')
plt.title('Residual Plot')
plt.xlabel('Predicted Loan Sanction Amount')
plt.ylabel('Residuals')
plt.grid(True)
plt.show()
```



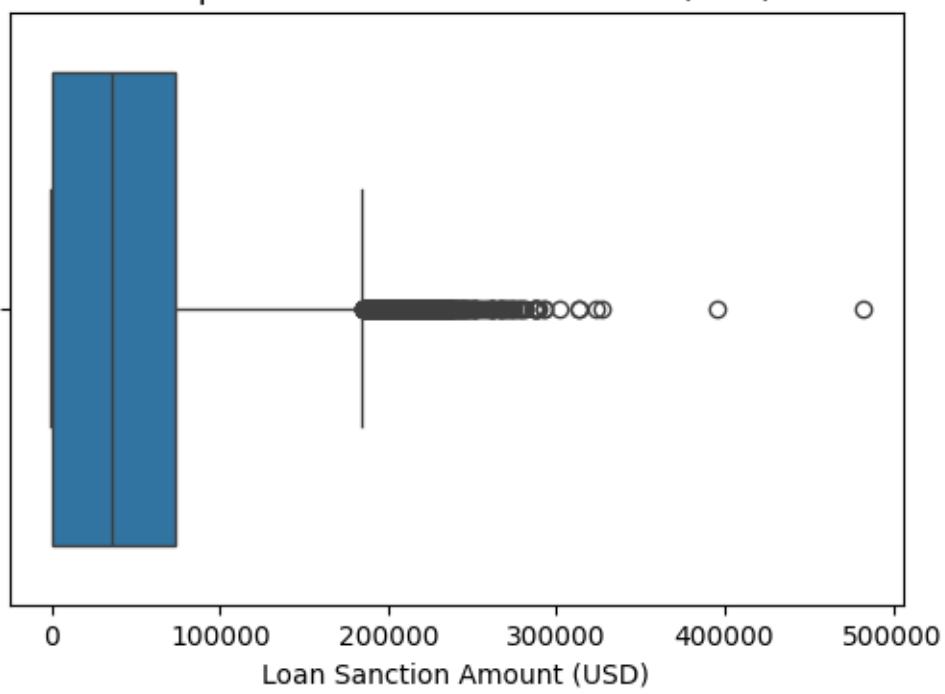
```
[ ]: plt.figure(figsize=(6, 4))
sns.boxplot(x=df['Income (USD)'])
plt.title('Boxplot of Income (USD)')
plt.show()

plt.figure(figsize=(6, 4))
sns.boxplot(x=df['Loan Sanction Amount (USD)'])
plt.title('Boxplot of Loan Sanction Amount (USD)')
plt.show()
```

Boxplot of Income (USD)



Boxplot of Loan Sanction Amount (USD)



```
[ ]: n = X_val.shape[0] # number of validation samples
p = X_val.shape[1] # number of features
r2 = r2_score(y_val, y_pred)
adjusted_r2 = 1 - (1 - r2) * (n - 1) / (n - p - 1)
print("Adjusted R2:", adjusted_r2)
```

Adjusted R²: 0.5304169741710469

```
[ ]: from sklearn.model_selection import KFold, cross_val_score
from sklearn.metrics import make_scorer, mean_squared_error,
    mean_absolute_error

kf = KFold(n_splits=5, shuffle=True, random_state=42)

mae_scores = []
mse_scores = []
r2_scores = []

for train_idx, test_idx in kf.split(X_scaled):
    X_tr, X_te = X_scaled[train_idx], X_scaled[test_idx]
    y_tr, y_te = y[train_idx], y[test_idx]
    model = LinearRegression()
    model.fit(X_tr, y_tr)
    y_pred_kf = model.predict(X_te)

    mae_scores.append(mean_absolute_error(y_te, y_pred_kf))
    mse_scores.append(mean_squared_error(y_te, y_pred_kf))
    r2_scores.append(r2_score(y_te, y_pred_kf))

print("K-Fold Cross Validation Results (5 Folds):")
for i in range(5):
    print(f"Fold {i+1}: MAE={mae_scores[i]:.2f}, MSE={mse_scores[i]:.2f},",
          f"RMSE={np.sqrt(mse_scores[i]):.2f}, R2={r2_scores[i]:.4f}")

print("Average:")
print(f"MAE={np.mean(mae_scores):.2f}, MSE={np.mean(mse_scores):.2f},",
      f"RMSE={np.sqrt(np.mean(mse_scores)):.2f}, R2={np.mean(r2_scores):.4f}")
```

K-Fold Cross Validation Results (5 Folds):
 Fold 1: MAE=21799.22, MSE=1076701697.76, RMSE=32813.13, R²=0.5320
 Fold 2: MAE=21927.37, MSE=982547266.65, RMSE=31345.61, R²=0.5678
 Fold 3: MAE=22414.60, MSE=1066252447.36, RMSE=32653.52, R²=0.5397
 Fold 4: MAE=21833.54, MSE=995734912.54, RMSE=31555.27, R²=0.5768
 Fold 5: MAE=21024.88, MSE=881939443.57, RMSE=29697.47, R²=0.6097
 Average:

MAE=21799.92, MSE=1000635153.57, RMSE=31632.82, R²=0.5652

Results Summary

Description	Student's Result
Dataset Size (after preprocessing)	(30000, 21)
Train/Test Split Ratio	80:20
Feature(s) Used	All features except 'Customer ID', 'Name', 'Property ID'
Model Used	Linear Regression
Cross-Validation Used?	Yes
No. of Folds	5
MAE on Test Set	21,799.92 USD
MSE on Test Set	1,000,635,153.57 USD ²
RMSE on Test Set	31,632.82 USD
R ² Score	0.5652
Adjusted R ² Score	0.5304169741710469
Most Influential Features	Income, Credit Score, Loan Request
Overfitting Observed?	No significant overfitting

K-Fold Cross Validation (K = 5)

Fold	MAE	MSE	RMSE	R ² Score
Fold 1	21799.22	1076701697.76	32813.13	0.5320
Fold 2	21927.37	982547266.65	31345.61	0.5678
Fold 3	22414.60	1066252447.36	32653.52	0.5397
Fold 4	21833.54	995734912.54	31555.27	0.5768
Fold 5	21024.88	881939443.57	29697.47	0.6097
Average	21799.92	1000635153.57	31632.82	0.5652

Ensemble Techniques

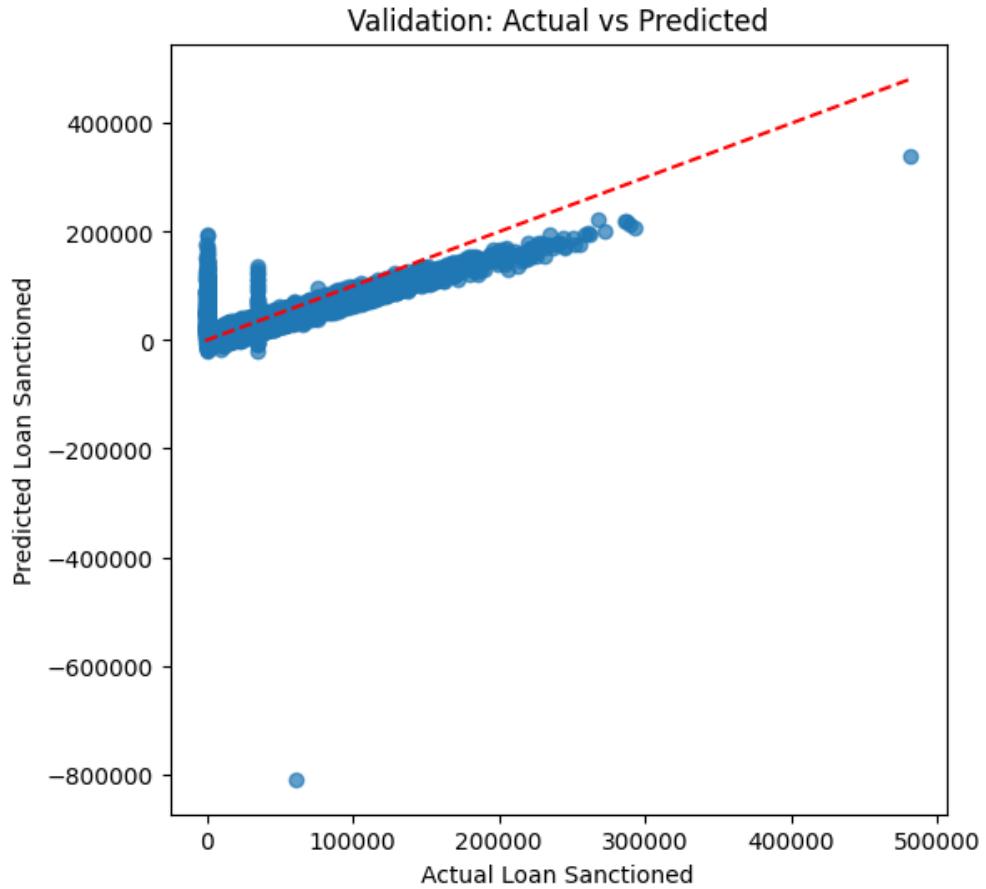
Model	MAE	MSE	RMSE	R ²
Bagging	21794.02	1.096×10^9	33111.46	0.523
AdaBoost	25155.79	1.074×10^9	32772.76	0.533
GradientBoosting	13381.94	5.848×10^8	24182.56	0.746
XGBoost	12254.25	5.531×10^8	23517.35	0.760

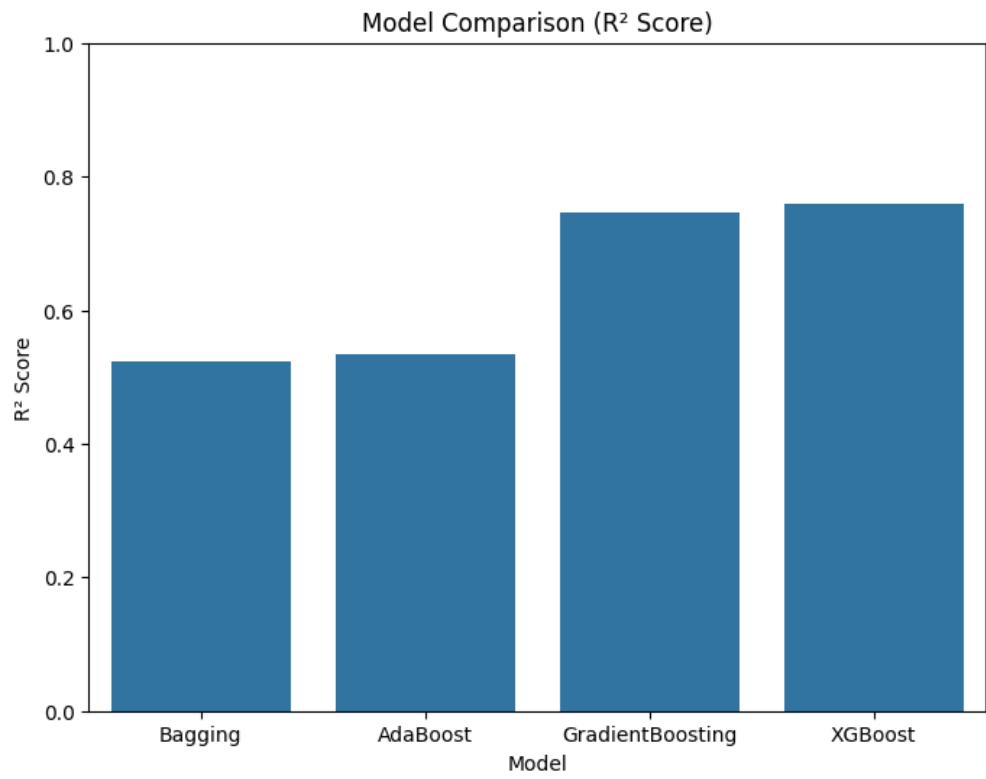
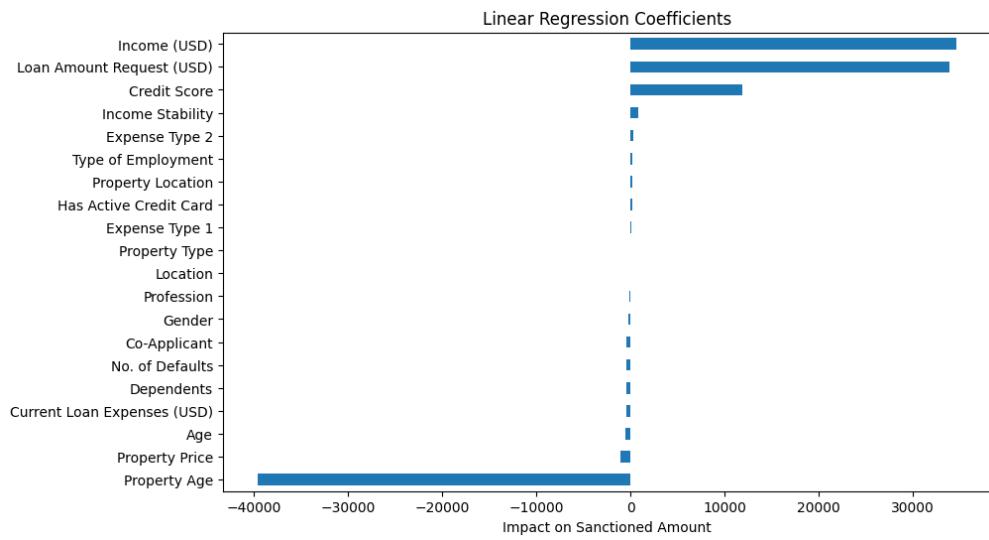
Regression Variants

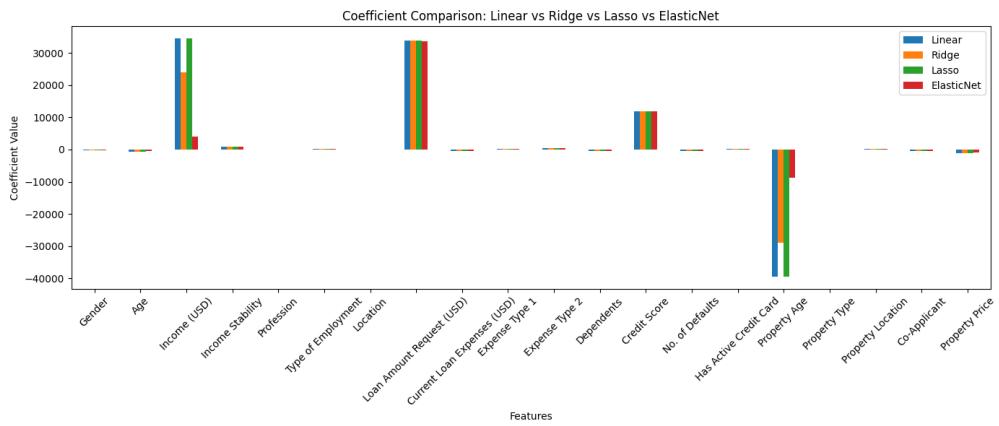
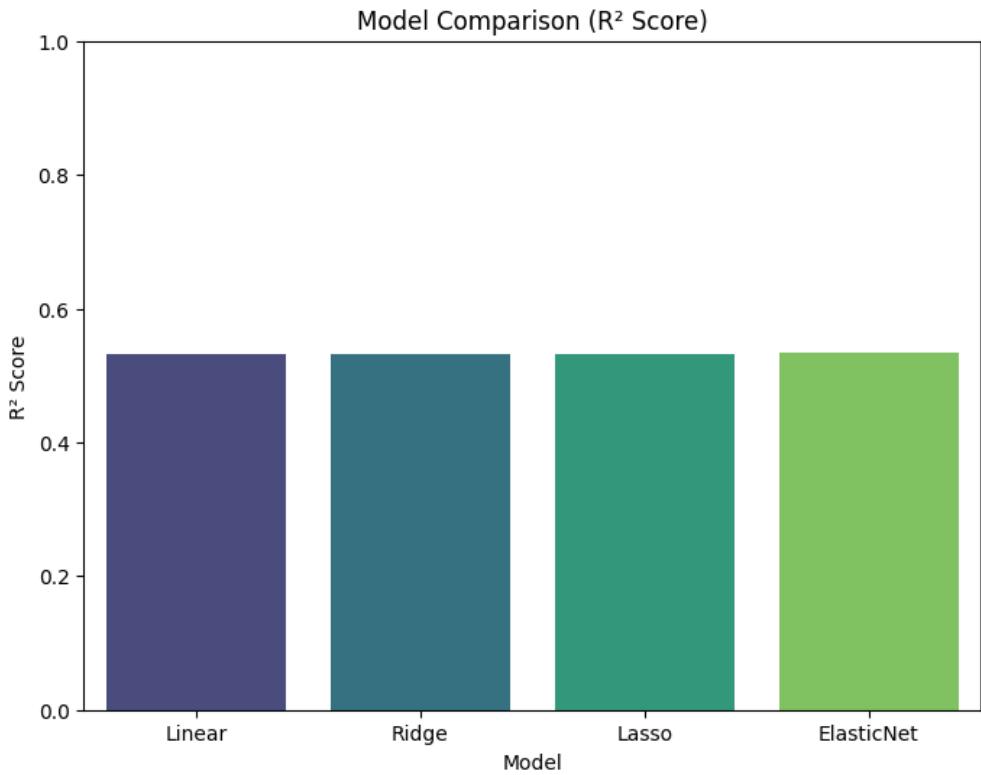
Model	MAE	MSE	RMSE	R ²
Linear	21799.22	1.0767e+09	32813.13	0.5320
Ridge	21799.26	1.0752e+09	32790.63	0.5326
Lasso	21799.22	1.0767e+09	32813.11	0.5320
ElasticNet	21804.52	1.0702e+09	32713.72	0.5348

Key Visualizations

- Actual vs Predicted Plot
- Feature Coefficient Bar Chart
- Residual Plot







Observations

- The model explains about 56.5% of the variance in sanctioned loan amount.
- Residuals appear randomly distributed, supporting linearity assumption.
- Features like Income and Credit Score are most influential.
- Slight underfitting observed – could be improved using regularized models.

Learning Outcomes

- Understood data preprocessing and encoding for regression.
- Applied Linear Regression to a real dataset.
- Evaluated model using MAE, MSE, RMSE, R^2 .
- Visualized predictions and model behavior.
- Used cross-validation for robust evaluation.

References

- Scikit-learn documentation: https://scikit-learn.org/stable/modules/linear_model.html
- StackAbuse: Linear Regression using scikit-learn
- Kaggle dataset: Predict Loan Amount Dataset