

# Spam Classification using Naïve Bayes, KNN, and SVM

Machine Learning Algorithms Laboratory (ICS1512)

September 3, 2025

## 1 Introduction

The objective of this experiment is to build and evaluate different machine learning models for classifying emails as *spam* or *ham*. The models compared are:

- Naïve Bayes Classifiers (GaussianNB, MultinomialNB, BernoulliNB)
- K-Nearest Neighbors (KNN) with varying hyperparameters
- Support Vector Machines (SVM) with different kernels

We evaluate models using accuracy, precision, recall, F1-score, confusion matrices, ROC curves, and cross-validation.

## 2 Dataset

The dataset used is the **Spambase dataset** [1], downloaded from Kaggle. It contains 4601 samples with 57 numerical features and a binary class label:

- 0 = Ham (not spam)
- 1 = Spam

## 3 Exploratory Data Analysis (EDA)

- The dataset contains no missing values.
- Class distribution was visualized using a count plot (Figure ??).

## 4 Data Preprocessing

- Features were standardized using `StandardScaler`.
- The dataset was split into training, validation, and test sets using stratified sampling.

## 5 Code and Output

```
[ ]: import kagglehub

# Download latest version
path = kagglehub.dataset_download("somesh24/spambase")

print("Path to dataset files:", path)
```

Path to dataset files: /kaggle/input/spambase

```
[ ]: import os

# List all files in the dataset folder
print(os.listdir(path))
```

['spambase\_csv.csv']

```
[ ]: import pandas as pd

df = pd.read_csv(os.path.join(path, 'spambase_csv.csv')) # adjust name if needed
print(df.head())
```

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	\
0	0.00	0.64	0.64	0.0	
1	0.21	0.28	0.50	0.0	
2	0.06	0.00	0.71	0.0	
3	0.00	0.00	0.00	0.0	
4	0.00	0.00	0.00	0.0	

	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet	\
0	0.32	0.00	0.00	0.00	
1	0.14	0.28	0.21	0.07	
2	1.23	0.19	0.19	0.12	
3	0.63	0.00	0.31	0.63	
4	0.63	0.00	0.31	0.63	

	word_freq_order	word_freq_mail	...	char_freq_%3B	char_freq_%28	\
0	0.00	0.00	...	0.00	0.000	
1	0.00	0.94	...	0.00	0.132	
2	0.64	0.25	...	0.01	0.143	
3	0.31	0.63	...	0.00	0.137	
4	0.31	0.63	...	0.00	0.135	

	char_freq_%5B	char_freq_%21	char_freq_%24	char_freq_%23	\
0	0.0	0.778	0.000	0.000	
1	0.0	0.372	0.180	0.048	
2	0.0	0.276	0.184	0.010	
3	0.0	0.137	0.000	0.000	

```

4          0.0      0.135      0.000      0.000
           capital_run_length_average  capital_run_length_longest \
0                  3.756                      61
1                  5.114                     101
2                  9.821                     485
3                  3.537                      40
4                  3.537                      40

           capital_run_length_total  class
0                  278          1
1                 1028          1
2                 2259          1
3                  191          1
4                  191          1

[5 rows x 58 columns]

```

```
[ ]: print(df.shape)
print(df.dtypes)
```

```
(4601, 58)
word_freq_make            float64
word_freq_address          float64
word_freq_all              float64
word_freq_3d               float64
word_freq_our              float64
word_freq_over             float64
word_freq_remove           float64
word_freq_internet         float64
word_freq_order            float64
word_freq_mail             float64
word_freq_receive          float64
word_freq_will              float64
word_freq_people            float64
word_freq_report            float64
word_freq_addresses         float64
word_freq_free              float64
word_freq_business           float64
word_freq_email             float64
word_freq_you               float64
word_freq_credit             float64
word_freq_your              float64
word_freq_font              float64
word_freq_000               float64
word_freq_money             float64
word_freq_hp                float64
word_freq_hpl               float64
```

```
word_freq_george          float64
word_freq_650              float64
word_freq_lab              float64
word_freq_labs             float64
word_freq_telnet            float64
word_freq_857              float64
word_freq_data              float64
word_freq_415              float64
word_freq_85                float64
word_freq_technology        float64
word_freq_1999              float64
word_freq_parts             float64
word_freq_pm                float64
word_freq_direct            float64
word_freq_cs                float64
word_freq_meeting            float64
word_freq_original           float64
word_freq_project            float64
word_freq_re                float64
word_freq_edu               float64
word_freq_table              float64
word_freq_conference         float64
char_freq_%3B               float64
char_freq_%28               float64
char_freq_%5B               float64
char_freq_%21               float64
char_freq_%24               float64
char_freq_%23               float64
capital_run_length_average float64
capital_run_length_longest   int64
capital_run_length_total    int64
class                      int64
dtype: object
```

##2. EDA


Split features and label

```
[ ]: X = df.drop('class', axis=1)
y = df['class']
```

Visualize Class Balance

```
[ ]: import seaborn as sns
import matplotlib.pyplot as plt

sns.countplot(x=y)
plt.title("Spam vs Ham Distribution (0 = Ham, 1 = Spam)")
plt.show()
```



Check for Missing Values

```
[ ]: print(X.isnull().sum().sum()) # Should be 0
```

0

##3. Scale the Features

```
[ ]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

## 5.1 4. Splitting (Train, Valid, Test)

```
[ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
```

##5. Model Building - Naive Bayes

```
[ ]: from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB

models = {
    "GaussianNB": GaussianNB(),
    "MultinomialNB": MultinomialNB(),
    "BernoulliNB": BernoulliNB()
}

[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
     f1_score, confusion_matrix, roc_curve, auc, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

results = []

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:, 1] if hasattr(model, u
     ↪'predict_proba') else None

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='binary')
    rec = recall_score(y_test, y_pred, average='binary')
    f1 = f1_score(y_test, y_pred, average='binary')

    results.append([name, acc, prec, rec, f1])

# Create subplots for CM and ROC
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# --- Confusion Matrix ---
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot(ax=axes[0], colorbar=False)
axes[0].set_title(f"Confusion Matrix - {name}")





# --- ROC Curve ---
if y_proba is not None:
    fpr, tpr, _ = roc_curve(y_test, y_proba)
    roc_auc = auc(fpr, tpr)
    axes[1].plot(fpr, tpr, label=f"AUC = {roc_auc:.2f}")
    axes[1].plot([0, 1], [0, 1], 'k--')
    axes[1].set_title(f"ROC Curve - {name}")
    axes[1].set_xlabel("False Positive Rate (FPR)")
    axes[1].set_ylabel("True Positive Rate (TPR)")
    axes[1].legend(loc="lower right")
    axes[1].grid(True)
```


```

else:
    axes[1].text(0.5, 0.5, "ROC Not Available", ha="center", va="center", u
    ↪fontsize=12)
    axes[1].set_axis_off()

plt.tight_layout()
plt.show()

```





```
[ ]: results_df = pd.DataFrame(results, columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1-Score'])
results_df = results_df.sort_values(by='F1-Score', ascending=False)
print("\nModel Performance Summary:\n")
print(results_df)
```

Model Performance Summary:

	Model	Accuracy	Precision	Recall	F1-Score
2	BernoulliNB	0.876221	0.871642	0.804408	0.836676
0	GaussianNB	0.833876	0.717842	0.953168	0.818935
1	MultinomialNB	0.776330	0.719888	0.707989	0.713889

##5.2 Model Building - KNN

```
[ ]: from sklearn.neighbors import KNeighborsClassifier

# Try multiple Ks and algorithms
models = {}
for k in [1, 3, 5, 7]:
    models[f"KNN_k={k}_kd"] = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree')
    models[f"KNN_k={k}_ball"] = KNeighborsClassifier(n_neighbors=k, algorithm='ball_tree')
```

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, roc_curve, auc, ConfusionMatrixDisplay
)
```

```

from sklearn.neighbors import KNeighborsClassifier
import pandas as pd
import matplotlib.pyplot as plt

# --- Step 1: Split into train/val/test ---
X_train_full, X_temp, y_train_full, y_temp = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp
)

# --- Stage 1: Vary k (algorithm fixed to kd_tree) ---
k_values = [1, 3, 5, 7]
k_results = []

for k in k_values:
    model = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree')
    model.fit(X_train_full, y_train_full)
    y_pred_val = model.predict(X_val)

    acc = accuracy_score(y_val, y_pred_val)
    prec = precision_score(y_val, y_pred_val, average='binary')
    rec = recall_score(y_val, y_pred_val, average='binary')
    f1 = f1_score(y_val, y_pred_val, average='binary')

    k_results.append([k, acc, prec, rec, f1])

df_k = pd.DataFrame(k_results, columns=["k", "Accuracy", "Precision", "Recall", "F1"])
print("\nPerformance for varying k (kd_tree):\n", df_k)

best_k = df_k.sort_values("Accuracy", ascending=False).iloc[0]["k"]
print(f"\nBest k value: {best_k}")

# --- Stage 2: Compare kd_tree vs ball_tree for best k ---
algorithms = ["kd_tree", "ball_tree"]
algo_results = []

for algo in algorithms:
    model = KNeighborsClassifier(n_neighbors=int(best_k), algorithm=algo)
    model.fit(X_train_full, y_train_full)
    y_pred_val = model.predict(X_val)

    acc = accuracy_score(y_val, y_pred_val)
    prec = precision_score(y_val, y_pred_val, average='binary')
    rec = recall_score(y_val, y_pred_val, average='binary')

```

```

f1 = f1_score(y_val, y_pred_val, average='binary')

algo_results.append([algo, acc, prec, rec, f1])

df_algo = pd.DataFrame(algo_results, columns=["Algorithm", "Accuracy", "Precision", "Recall", "F1"])
print("\nPerformance for kd_tree vs ball_tree:\n", df_algo)

best_algo = df_algo.sort_values("Accuracy", ascending=False).iloc[0]["Algorithm"]
print(f"\nBest algorithm: {best_algo}")

# --- Stage 3: Final evaluation on test set ---
final_model = KNeighborsClassifier(n_neighbors=int(best_k), algorithm=best_algo)
final_model.fit(X_train_full, y_train_full)
y_pred_test = final_model.predict(X_test)
y_proba_test = final_model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred_test)
prec = precision_score(y_test, y_pred_test, average='binary')
rec = recall_score(y_test, y_pred_test, average='binary')
f1 = f1_score(y_test, y_pred_test, average='binary')

print("\nFinal Test Set Performance:")
print(f"Accuracy: {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall: {rec:.4f}")
print(f"F1 Score: {f1:.4f}")

# --- Plots ---
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred_test)
disp = ConfusionMatrixDisplay(cm)
disp.plot(ax=axes[0], colorbar=False)
axes[0].set_title(f"Confusion Matrix - K={best_k}, Algo={best_algo}")

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_proba_test)
roc_auc = auc(fpr, tpr)
axes[1].plot(fpr, tpr, label=f"AUC = {roc_auc:.2f}")
axes[1].plot([0, 1], [0, 1], 'k--')
axes[1].set_title(f"ROC Curve - K={best_k}, Algo={best_algo}")
axes[1].set_xlabel("False Positive Rate (FPR)")
axes[1].set_ylabel("True Positive Rate (TPR)")
axes[1].legend(loc="lower right")
axes[1].grid(True)

```

```
plt.tight_layout()  
plt.show()
```

Performance for varying k (kd\_tree):

k	Accuracy	Precision	Recall	F1	
0	1	0.807246	0.745583	0.775735	0.760360
1	3	0.797101	0.742647	0.742647	0.742647
2	5	0.797101	0.750000	0.727941	0.738806
3	7	0.804348	0.758491	0.738971	0.748603

Best k value: 1.0

Performance for kd\_tree vs ball\_tree:

	Algorithm	Accuracy	Precision	Recall	F1
0	kd_tree	0.807246	0.745583	0.775735	0.76036
1	ball_tree	0.807246	0.745583	0.775735	0.76036

Best algorithm: kd\_tree



Final Test Set Performance:

Accuracy: 0.7873

Precision: 0.7178

Recall: 0.7574

F1 Score: 0.7370



```
[ ]: import time  
from sklearn.svm import SVC  
from sklearn.model_selection import cross_val_score, StratifiedKFold
```

```

# =====
# SVM Kernel-wise Evaluation
# =====
svm_results = []

kernels = {
    "linear": {"C": [1]},
    "poly": {"C": [1], "degree": [2], "gamma": ["scale"]},
    "rbf": {"C": [1], "gamma": ["scale"]},
    "sigmoid": {"C": [1], "gamma": ["scale"]}}
}

best_svm = None
best_score = 0

for kernel, params in kernels.items():
    for C in params.get("C", [1]):
        for degree in params.get("degree", [3]):
            for gamma in params.get("gamma", ["scale"]):
                start = time.time()
                if kernel == "linear":
                    model = SVC(kernel=kernel, C=C, probability=True,
                     random_state=42)
                elif kernel == "poly":
                    model = SVC(kernel=kernel, C=C, degree=degree, gamma=gamma,
                     probability=True, random_state=42)
                else:
                    model = SVC(kernel=kernel, C=C, gamma=gamma,
                     probability=True, random_state=42)

                model.fit(X_train, y_train)
                train_time = time.time() - start
                y_pred = model.predict(X_test)

                acc = accuracy_score(y_test, y_pred)
                f1 = f1_score(y_test, y_pred)

                svm_results.append([kernel, f"C={C}", degree=degree if
                 kernel=='poly' else '-'], gamma=gamma if kernel!='linear' else '-',
                 acc, f1, train_time])

                if acc > best_score:
                    best_score = acc
                    best_svm = model

```

```

svm_table = pd.DataFrame(svm_results, columns=["Kernel", "Hyperparameters", "Accuracy", "F1 Score", "Training Time"])
print("\n Table 4: SVM Performance with Different Kernels")
display(svm_table)

```

Table 4: SVM Performance with Different Kernels

	Kernel	Hyperparameters	Accuracy	F1 Score	Training Time
0	linear	C=1, degree=-, gamma=-	0.926194	0.903226	2088.215490
1	poly	C=1, degree=2, gamma=scale	0.668596	0.343840	5.502092
2	rbf	C=1, degree=-, gamma=scale	0.726483	0.590022	5.392869
3	sigmoid	C=1, degree=-, gamma=scale	0.589001	0.483636	4.809869

```

[ ]: # =====
# 5-Fold Cross Validation
# =====
import numpy as np
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Use your best models
nb_model = GaussianNB() # Best NB
knn_model = final_model # Best KNN found earlier
svm_model = best_svm # Best SVM from kernel search

models_cv = {"Naïve Bayes": nb_model, "KNN": knn_model, "SVM": svm_model}

cv_results = {"Fold": [], "Naïve Bayes Accuracy": [], "KNN Accuracy": [], "SVM Accuracy": []}

# Use scaled features for fairness
for fold, (train_idx, val_idx) in enumerate(cv.split(X_scaled, y), 1):
    X_tr, X_val = X_scaled[train_idx], X_scaled[val_idx]
    y_tr, y_val = y.values[train_idx], y.values[val_idx]

    cv_results["Fold"].append(f"Fold {fold}")
    for name, model in models_cv.items():
        model.fit(X_tr, y_tr)
        y_pred = model.predict(X_val)
        acc = accuracy_score(y_val, y_pred)
        cv_results[f"{name} Accuracy"].append(acc)

# Add averages
cv_results["Fold"].append("Average")
for name in models_cv.keys():

```

```

avg = np.mean(cv_results[f"{name} Accuracy"])
cv_results[f"{name} Accuracy"].append(avg)

cv_table = pd.DataFrame(cv_results)
print("\n Cross-Validation Accuracy Scores:\n")
display(cv_table)

```

Cross-Validation Accuracy Scores:

	Fold	Naïve Bayes Accuracy	KNN Accuracy	SVM Accuracy
0	Fold 1	0.820847	0.909881	0.917481
1	Fold 2	0.808696	0.907609	0.934783
2	Fold 3	0.801087	0.922826	0.927174
3	Fold 4	0.821739	0.910870	0.927174
4	Fold 5	0.825000	0.908696	0.923913
5	Average	0.815474	0.911976	0.926105

[2] :

## 6 Model Building and Results

### 6.1 Naïve Bayes Classifiers

Three variants were trained and evaluated: GaussianNB, MultinomialNB, and BernoulliNB. Table 1 summarizes the results.

Model	Accuracy	Precision	Recall	F1-Score
GaussianNB	0.8339	0.7178	0.9532	0.8189
MultinomialNB	0.7763	0.7199	0.7080	0.7139
BernoulliNB	0.8762	0.8716	0.8044	0.8367

### 6.2 K-Nearest Neighbors (KNN)

k	Accuracy	Precision	Recall	F1
1	0.8072	0.7456	0.7757	0.7604
3	0.7971	0.7426	0.7426	0.7426
5	0.7971	0.7500	0.7279	0.7388
7	0.8043	0.7585	0.7390	0.7486

Best  $k$  value: 1.0

Algorithm	Accuracy	Precision	Recall	F1
kd_tree	0.8072	0.7456	0.7757	0.7604
ball_tree	0.8072	0.7456	0.7757	0.7604

Best algorithm: kd \_ tree

Final Test Set Performance:

- Accuracy: 0.7873
- Precision: 0.7178
- Recall: 0.7574
- F1 Score: 0.7370

### 6.3 Support Vector Machines (SVM)

Kernel	Hyperparameters	Accuracy	F1-Score	Training Time (seconds)
Linear	C=1	0.9262	0.9032	2088.22
Polynomial	C=1, degree=2, gamma=scale	0.6686	0.3438	5.50
RBF	C=1, gamma=scale	0.7265	0.5900	5.39
Sigmoid	C=1, gamma=scale	0.5890	0.4836	4.81

## 7 Cross-Validation

To ensure robustness, 5-fold stratified cross-validation was conducted on the best-performing models from each category. Table 5 shows fold-wise and average accuracies.

Fold	Naïve Bayes Accuracy	KNN Accuracy	SVM Accuracy
Fold 1	0.8208	0.9099	0.9175
Fold 2	0.8087	0.9076	0.9348
Fold 3	0.8011	0.9228	0.9272
Fold 4	0.8217	0.9109	0.9272
Fold 5	0.8250	0.9087	0.9239
Average	0.8155	0.9120	0.9261

## 8 Conclusion

- Naïve Bayes provided a strong baseline with fast training time.
- KNN performance was highly dependent on  $k$  and algorithm choice.
- SVM with RBF kernel gave the most balanced results, achieving the highest accuracy and F1-score.

## References

- [1] Spambase Dataset. Available at: <https://www.kaggle.com/datasets/somesh24/spambase>