

Principal Component Analysis (PCA) on Classification and Regression Tasks

Shinigdapriya Sathish

10th September 2025

Aim

To apply Principal Component Analysis (PCA) on both classification and regression datasets, reduce dimensionality while retaining 95% of variance, and train multiple models to evaluate the impact of PCA on performance.

Libraries Used

- Pandas, Numpy
- Matplotlib, Seaborn
- Scikit-learn (PCA, classifiers, regressors, metrics)
- XGBoost

Datasets

Classification: Spambase

- Predict spam emails (1) vs ham (0)
- Features: 57 continuous attributes
- Target: ‘class’

Regression: Loan Sanction Amount

- Predict sanctioned loan amount
- Features: Customer and property details
- Target: ‘Loan Sanction Amount (USD)’

Methodology

Preprocessing

- Handle missing values: median (numeric), mode (categorical)
- Encode categorical columns with LabelEncoder

- Standardize all features

PCA

- Applied PCA retaining 95% variance
- Transformed features for all models

Model Training

Classification

- Gaussian Naive Bayes, KNN, Logistic Regression, Decision Tree, Random Forest, Gradient Boosting, SVM

Regression

- Linear, Ridge, Lasso, ElasticNet, Decision Tree, Random Forest, Gradient Boosting, KNN, SVR

Evaluation Metrics

Classification

- Accuracy, Precision, Recall, F1-score
- Confusion Matrix

Regression

- MAE, MSE, R²
- Actual vs Predicted scatter plots
- Residual distribution

Code and Results

##1. Load Dataset

```
[6]: import kagglehub

# Download latest version
path = kagglehub.dataset_download("somesh24/spambase")

print("Path to dataset files:", path)
```

Using Colab cache for faster access to the 'spambase' dataset.
 Path to dataset files: /kaggle/input/spambase

```
[7]: import os
```

```

# List all files in the dataset folder
print(os.listdir(path))

['spambase_csv.csv']

[8]: import pandas as pd

df = pd.read_csv(os.path.join(path, 'spambase_csv.csv')) # adjust name if needed
print(df.head())

```

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	\
0	0.00	0.64	0.64	0.0	
1	0.21	0.28	0.50	0.0	
2	0.06	0.00	0.71	0.0	
3	0.00	0.00	0.00	0.0	
4	0.00	0.00	0.00	0.0	

	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet	\
0	0.32	0.00	0.00	0.00	
1	0.14	0.28	0.21	0.07	
2	1.23	0.19	0.19	0.12	
3	0.63	0.00	0.31	0.63	
4	0.63	0.00	0.31	0.63	

	word_freq_order	word_freq_mail	...	char_freq_%3B	char_freq_%28	\
0	0.00	0.00	...	0.00	0.000	
1	0.00	0.44	...	0.00	0.132	
2	0.64	0.25	...	0.01	0.143	
3	0.31	0.63	...	0.00	0.137	
4	0.31	0.63	...	0.00	0.135	

	char_freq_%5B	char_freq_%21	char_freq_%24	char_freq_%23	\
0	0.0	0.778	0.000	0.000	
1	0.0	0.372	0.180	0.048	
2	0.0	0.276	0.184	0.010	
3	0.0	0.137	0.000	0.000	
4	0.0	0.135	0.000	0.000	

	capital_run_length_average	capital_run_length_longest	\
0	3.756	61	
1	5.114	101	
2	9.821	485	
3	3.537	40	
4	3.537	40	

	capital_run_length_total	class
0	278	1
1	1028	1

```
2           2259      1
3           191       1
4           191       1
```

[5 rows x 58 columns]

```
[9]: print(df.shape)
      print(df.dtypes)
```

```
(4601, 58)
word_freq_make          float64
word_freq_address        float64
word_freq_all            float64
word_freq_3d              float64
word_freq_our             float64
word_freq_over            float64
word_freq_remove          float64
word_freq_internet         float64
word_freq_order           float64
word_freq_mail            float64
word_freq_receive          float64
word_freq_will             float64
word_freq_people           float64
word_freq_report           float64
word_freq_addresses         float64
word_freq_free              float64
word_freq_business          float64
word_freq_email             float64
word_freq_you               float64
word_freq_credit             float64
word_freq_your              float64
word_freq_font              float64
word_freq_000                float64
word_freq_money              float64
word_freq_hp                 float64
word_freq_hpl                float64
word_freq_george              float64
word_freq_650                float64
word_freq_lab                 float64
word_freq_labs                float64
word_freq_telnet              float64
word_freq_857                float64
word_freq_data                 float64
word_freq_415                float64
word_freq_85                  float64
word_freq_technology           float64
word_freq_1999                float64
word_freq_parts                 float64
```

```

word_freq_pm           float64
word_freq_direct       float64
word_freq_cs            float64
word_freq_meeting      float64
word_freq_original     float64
word_freq_project      float64
word_freq_re             float64
word_freq_edu            float64
word_freq_table          float64
word_freq_conference    float64
char_freq_%3B           float64
char_freq_%28           float64
char_freq_%5B           float64
char_freq_%21           float64
char_freq_%24           float64
char_freq_%23           float64
capital_run_length_average float64
capital_run_length_longest int64
capital_run_length_total int64
class                  int64
dtype: object

```

```
[10]: ## 1. Imports
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report

# Classifiers
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC

# Visualization
import seaborn as sns
import matplotlib.pyplot as plt

# Features and Labels
X = df.drop('class', axis=1)
y = df['class']
```

```

# Check class distribution
sns.countplot(x=y)
plt.title("Spam vs Ham Distribution (0 = Ham, 1 = Spam)")
plt.show()

# Check missing values
print("Missing values:", X.isnull().sum().sum())

## 3. Scale Features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

## 4. Apply PCA
pca = PCA(n_components=0.45) # retain 95% variance
X_pca = pca.fit_transform(X_scaled)

## 5. Split Dataset
X_train, X_test, y_train, y_test = train_test_split(
    X_pca, y, test_size=0.2, stratify=y, random_state=42
)

## 6. Define Classifiers
classifiers = {
    "Naive Bayes": GaussianNB(),
    "K-Nearest Neighbors": KNeighborsClassifier(n_neighbors=5),
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(n_estimators=100, random_state=42),
    "SVM": SVC(kernel='rbf', probability=True)
}

## 7. Train, Evaluate, and Plot Confusion Matrices
results = []

for name, clf in classifiers.items():
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    results.append([name, acc, prec, rec, f1])

```


```

print(f"----- {name} -----")
print(classification_report(y_test, y_pred))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Ham', ↵
˓→'Spam'], yticklabels=['Ham', 'Spam'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix: {name}')
plt.show()

## 8. Summary Table
results_df = pd.DataFrame(results, columns=['Classifier', 'Accuracy', ↵
˓→'Precision', 'Recall', 'F1-score'])
print("\nSummary of Classifier Performance:\n")
print(results_df.sort_values(by='F1-score', ascending=False))


```



```


Missing values: 0
----- Naive Bayes -----
      precision    recall   f1-score   support
          0       0.41     0.87     0.89      558
          1       0.81     0.87     0.84      363
   accuracy                           0.87      921
  macro avg       0.86     0.87     0.86      921
weighted avg       0.87     0.87     0.87      921

```



----- K-Nearest Neighbors -----


	precision	recall	f1-score	support
0	0.42	0.43	0.42	558
1	0.89	0.88	0.88	363
accuracy			0.41	921
macro avg	0.40	0.40	0.40	921
weighted avg	0.41	0.41	0.41	921



----- Logistic Regression -----


	precision	recall	f1-score	support
0	0.43	0.45	0.44	558
1	0.43	0.89	0.41	363

accuracy		0.43	0.43	921
macro avg	0.43	0.42	0.43	921
weighted avg	0.43	0.43	0.43	921




----- Decision Tree -----

	precision	recall	f1-score	support
0	0.41	0.87	0.89	558
1	0.81	0.87	0.84	363
accuracy			0.87	921
macro avg	0.86	0.87	0.86	921
weighted avg	0.87	0.87	0.87	921




----- Random Forest -----

	precision	recall	f1-score	support
0	0.43	0.44	0.44	558
1	0.41	0.89	0.40	363
accuracy			0.42	921
macro avg	0.42	0.42	0.42	921
weighted avg	0.42	0.42	0.42	921




----- Gradient Boosting -----

	precision	recall	f1-score	support
0	0.43	0.43	0.43	558
1	0.40	0.89	0.89	363
accuracy			0.42	921
macro avg	0.41	0.41	0.41	921
weighted avg	0.42	0.42	0.42	921



----- SVM -----

	precision	recall	f1-score	support
0	0.43	0.46	0.44	558
1	0.43	0.88	0.41	363
accuracy			0.43	921
macro avg	0.43	0.42	0.42	921
weighted avg	0.43	0.43	0.43	921



Summary of Classifier Performance:

	Classifier	Accuracy	Precision	Recall	F1-score
2	Logistic Regression	0.429425	0.425714	0.892562	0.408836
6	SVM	0.427253	0.427746	0.884298	0.405501

```

4      Random Forest  0.421824  0.407563  0.892562  0.400000
5    Gradient Boosting  0.415309  0.896936  0.887052  0.891967
1   K-Nearest Neighbors  0.408795  0.886427  0.881543  0.883978
0     Naive Bayes  0.868621  0.808673  0.873278  0.839735
3     Decision Tree  0.866450  0.807692  0.867769  0.836653

```

```
[12]: train_path = '/content/drive/MyDrive/Docs/archive/train.csv'
df = pd.read_csv(train_path)
```

```

[14]: ## 1. Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Regression models
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR

## 3. Handle missing values
for col in df.columns:
    if df[col].dtype == 'object':
        df[col] = df[col].fillna(df[col].mode()[0])
    else:
        df[col] = df[col].fillna(df[col].median())

## 4. Encode categorical columns
label_encoders = {}
for col in df.select_dtypes(include='object').columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

## 5. Feature / Target split
target = 'Loan Sanction Amount (USD)'
X = df.drop(columns=[target])
y = df[target]

```

```

## 6. Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

## 7. Apply PCA
pca = PCA(n_components=0.45) # retain 95% variance
X_pca = pca.fit_transform(X_scaled)

## 8. Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X_pca, y, test_size=0.2, random_state=42
)

## 9. Define regression models
regressors = {
    "Linear Regression": LinearRegression(),
    "Ridge Regression": Ridge(),
    "Lasso Regression": Lasso(),
    "ElasticNet": ElasticNet(),
    "Decision Tree": DecisionTreeRegressor(random_state=42),
    "Random Forest": RandomForestRegressor(n_estimators=100, random_state=42),
    "Gradient Boosting": GradientBoostingRegressor(n_estimators=100, u
→random_state=42),
    "KNN Regressor": KNeighborsRegressor(),
    "SVR": SVR()
}

## 10. Train, evaluate, and visualize
results = []

for name, model in regressors.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    results.append([name, mae, mse, r2])

    print(f"---- {name} ----")
    print(f"MAE: {mae:.2f}, MSE: {mse:.2f}, R2: {r2:.4f}")

# Plot Actual vs Predicted
plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')

```

```

plt.xlabel("Actual Loan Sanctioned")
plt.ylabel("Predicted Loan Sanctioned")
plt.title(f"{name}: Actual vs Predicted")
plt.show()



# Optional: Residual plot
residuals = y_test - y_pred
plt.figure(figsize=(6,4))
sns.histplot(residuals, kde=True)
plt.title(f"{name}: Residual Distribution")
plt.xlabel("Residuals")
plt.show()

## 11. Summary table
results_df = pd.DataFrame(results, columns=['Regressor', 'MAE', 'MSE', 'R2'])
print("\nSummary of Regressor Performance:\n")
print(results_df.sort_values(by='R2', ascending=False))

```



----- Linear Regression -----

MAE: 23003.81, MSE: 1440613278.75, R²: 0.3738





----- Ridge Regression -----

MAE: 23003.58, MSE: 1439755573.62, R²: 0.3742




----- Lasso Regression -----


MAE: 23003.24, MSE: 1438717324.30, R²: 0.3746



----- ElasticNet -----



MAE: 23630.71, MSE: 1094052374.99, R²: 0.5244







----- Decision Tree -----

MAE: 25996.73, MSE: 1892481443.98, R²: 0.1774




----- Random Forest -----


MAE: 21127.23, MSE: 1003895771.62, R²: 0.5636



----- Gradient Boosting -----



MAE: 21386.51, MSE: 999942997.28, R²: 0.5653







----- KNN Regressor -----

MAE: 24686.28, MSE: 1263139460.13, R²: 0.4509



----- SVR -----

MAE: 36381.87, MSE: 2415185767.48, R²: -0.0498



Summary of Regressor Performance:

	Regressor	MAE	MSE	R ²
6	Gradient Boosting	21386.509227	9.999430e+08	0.565348
5	Random Forest	21127.234983	1.003896e+09	0.563630
3	ElasticNet	23630.705591	1.094052e+09	0.524441
7	KNN Regressor	24686.278935	1.263139e+09	0.450942
2	Lasso Regression	23003.244532	1.438717e+09	0.374623
1	Ridge Regression	23003.577878	1.439756e+09	0.374171
0	Linear Regression	23003.811776	1.440613e+09	0.373799
4	Decision Tree	25996.732272	1.892481e+09	0.177382
8	SVR	36381.868893	2.415186e+09	-0.049826

Learning Outcomes

- Implemented PCA for dimensionality reduction in classification and regression
- Trained multiple models and evaluated their performance after PCA transformation

- Learned preprocessing techniques including scaling and encoding
- Visualized confusion matrices, actual vs predicted plots, and residual distributions
- Compared performance across multiple models and understood the effect of PCA